

Python for MATLAB Users

Promoting Open Source Computer Vision Research

Kitware, Inc. / Google Research

June 2012



Matt
Leotta



Amitha
Perera



Patrick
Reynolds



Eran
Squires



Yong
Zhao



Varun
Ganapathi



This presentation is copyrighted by
Kitware, Inc.

distributed under the
Creative Commons by Attribution License 3.0
<http://creativecommons.org/licenses/by/3.0>

Source code used in this presentation and
in accompanying examples and exercises
is distributed under the
Apache License 2.0
<http://www.apache.org/licenses/LICENSE-2.0>

Approved for Public Release, Distribution Unlimited

What this tutorial is about

- Promote open source computer vision research.
 - Reduce the community's dependence on MATLAB.
 - Show MATLAB users how to migrate to Python.
 - Present a Python IDE similar to that of MATLAB.
- Introduce Python for scientific computing.
- Explain how to access functionality in Python from third party vision libraries (OpenCV, ITK).
- Cover some advanced computing topics using Python:
 - parallel processing with threads
 - GPU processing with PyCUDA and PyOpenCL
 - calling C/C++ code from Python
 - calling legacy MATLAB code from Python
- Provided examples and hands-on exercises for all of the above.
- Raffle of an ASUS Transformer Pad and more.

Schedule (Approximate)

8:30	Overview, Virtual Machine Setup
9:00	Intro to Python
10:00–10:30	Coffee Break
10:30	Python Scripting
11:15	Scientific Python – Numpy, Matplotlib, SciPy
12:30–1:30	Lunch
1:30	Using OpenCV
2:15	Using Simple ITK
3:00–3:30	Coffee Break
3:30	Parallel Processing / GPU Processing
4:15	Interfacing Python with C/C++
4:45	Interfacing Python with MATLAB
5:00	Wrap-up, Raffle

Virtual Machine Distribution

Virtual Machine Distribution

- Get a USB Memory Stick
- Install VirtualBox from it
- Import the virtual machine file
- Boot the virtual machine
- Get familiar with the environment and directories

Media Content

Directories and Files

- VirtualBoxInstallers

- VirtualBox-4.1.16-78094-Linux_amd64.run (64-bit Linux)
- VirtualBox-4.1.16-78094-Linux_x86.run (32-bit Linux)
- VirtualBox-4.1.16-78094-OSX.dmg (Mac)
- VirtualBox-4.1.16-78094-Win.exe (Windows)
- LinuxPackages
 - RPM (Fedora 14–17, Mandriva 2010–2011, RHEL/CentOS 4–6, SUSE 10–11, OpenSUSE 11.3–11.4)
 - DEB (Debian 5–7, Ubuntu 8.04, 10.04–12.04)

- VirtualMachine

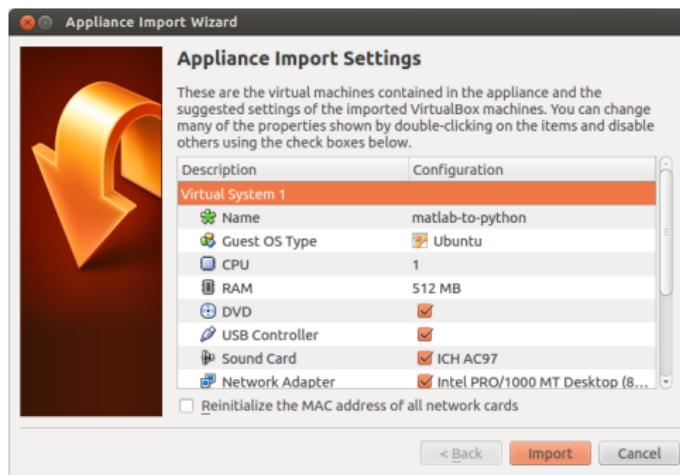
- matlab-to-python.ova

Install VirtualBox

- Select the installer for your platform
- Run it

Importing the Virtual Machine

- Run VirtualBox
- In “File” Menu select “Import Appliance”
- Provide the filename in the USB stick
“VirtualMachine/matlab-to-python.ova”
- You should see:



To Be Continued...

Why use Python?

- Sharing code and data is often overlooked but is as important as sharing algorithms and results in published scientific research.
- Shared code promotes reproducible research and scientific validation.
- Computer vision research code is commonly written in MATLAB.
- Sharing research code written in MATLAB is a good first step, but . . .
- MATLAB code can only be executed by a privileged few who can afford MATLAB license fees.

Why not MATLAB?

- Language design
 - MATLAB is effective for quick code and/or small code bases, but does not scale as projects grow
 - Numerical computing is easy, but more general programming is not
- Top-level functions and code organization
 - Function-name = file-name in MATLAB: discourages writing modular programs
 - Hard to appropriately name functions and organize code without conflicts
- Toolboxes are purchased separately
 - Shared MATLAB code will only run if you've licensed the right toolboxes
- Cost – Very Expensive!

MATLAB Pricing

MATLAB Component	Price
Individual Matlab License	\$2,150
MATLAB Compiler	\$5,000
Toolboxes	
Computer Vision System	\$1,350
Image Processing	\$1,000
Optimization	\$1,000
Parallel Computing	\$1,000
Statistics	\$1,000
Neural Network	\$1,000
Signal Processing	\$1,000
Symbolic Math	\$1,000
Total	\$15,500 <i>per user</i>

- Pricing as of April 12th, 2012.

More Info: <http://www.mathworks.com/store/productIndexLink.do>

Open Source Alternatives to MATLAB

- Octave (GPL license)
 - A mostly MATLAB compatible language
 - Command line only (no GUI IDE)
- FreeMat (GPL license)
 - A somewhat MATLAB compatible language
 - Very basic GUI
- Scilab (CeCILL license)
 - Not compatible with MATLAB language, provides a converter
 - Most full featured IDE
- These options solve the cost and license restrictions, but have the same technical limitations MATLAB.
- None of these options are perfectly compatible with MATLAB code.
- Each has a relatively small user community.

Why Python?

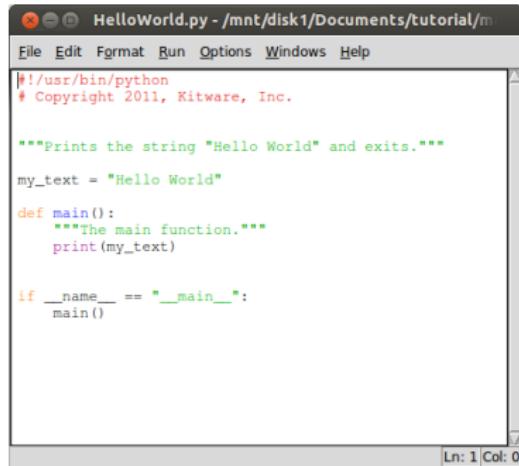
- Cost / License
 - Python and related packages are free (have no licensing fee)
 - Python and related packages use permissive BSD-like licenses.
- Language design
 - Python has well-designed language features supporting small and large code bases
 - Python code is highly readable
- Top-level functions and code organization
 - Python is modular and supports name spaces to avoid naming conflicts.
- Integration with other languages/tools
- An enormous user base.
 - Python is used for just about everything (web, database, networking, GUI, gaming, graphics, science, ...)
- Just search the web for testimonials

Ways to Use Python: Many IDE Choices

IDE	Developer	Latest stable release version	Latest stable release date	Platform	Toolkit	License
Anjuta	Naba Kumar	3.0.3.0	2011-06-10	Unix-like	GTK+	GPL
Eric Python IDE	Detlev Offenbach	4.4.18 and 5.1.6	2011-10-02	Independent	Qt	GPL
Geany	Team	0.20	2011-01-07	Independent	GTK2	GPL
IDLE	Guido van Rossum et al.	3.2	2011-05-13	Independent	Tkinter	Open source
Komodo IDE	ActiveState	6.0.0	2010-10-07	Cross-platform	Mozilla platform	Proprietary
MonoDevelop	Novell and the Mono community	2.4.2	2011-01-19	Cross-platform	Gtk#	GPL
NetBeans (7.0 and up will no longer support Python)	Oracle	6.9	2010-06-15	Cross-platform	Swing	GPL
Ninja IDE	Team	2.0-Beta3		Independent	Qt	GPL
PIDA	Team	0.6.2	2010-08-04	Cross-platform	PyGTK	GPL
PyCharm	JetBrains	2.0.2	2012-02-07	Cross-platform	Swing	Proprietary
PyDev	Aptana	2.4.0	2012-02-01	Eclipse (Cross-platform)	SWT	EPL
PyScripter		2.4.3	2011-09-20	Windows		MIT Licence
Python Tools for Visual Studio	Microsoft	1.1	2012-02-10	Windows		Apache License 2.0
Spyder	Pierre Raybaut et al.	2.1.9	2012-03-31	Independent	Qt	MIT Licence
Stanis Python Editor	Stani	0.8.4h	2008-02-14	Independent	wxPython	GPL
Wing	Wingware	4.1.3-1	2012-01-11	Linux/Windows/OS X	GTK2	Proprietary
wxGlade	Alberto Griggio	0.6.3	2008-02-02	Independent	wxPython	MIT License

- There are several choices for Python development environments.
- The above list is from Wikipedia (taken April 2012).
- Spyder is most like the MATLAB IDE.

Ways to Use Python: IDLE Python IDE



The screenshot shows the IDLE Python IDE interface. The title bar reads "HelloWorld.py - /mnt/disk1/Documents/tutorial/m". The menu bar includes File, Edit, Format, Run, Options, Windows, and Help. The main window displays the Python code for "HelloWorld.py". The code prints "Hello World" when run.

```
#!/usr/bin/python
# Copyright 2011, Kitware, Inc.

"""Prints the string "Hello World" and exits."""

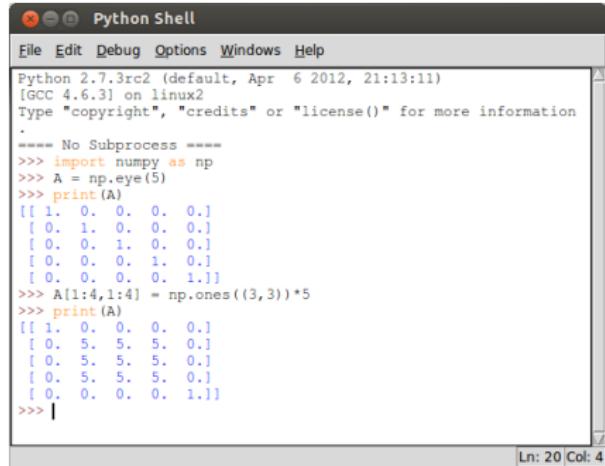
my_text = "Hello World"

def main():
    """The main function."""
    print(my_text)

if __name__ == "__main__":
    main()


```

Ln: 1 Col: 0



The screenshot shows the Python Shell window. The title bar reads "Python Shell". The menu bar includes File, Edit, Debug, Options, Windows, and Help. The shell window displays a session where a 5x5 identity matrix is created and multiplied by a 3x3 matrix of ones, resulting in a 5x3 matrix of ones.

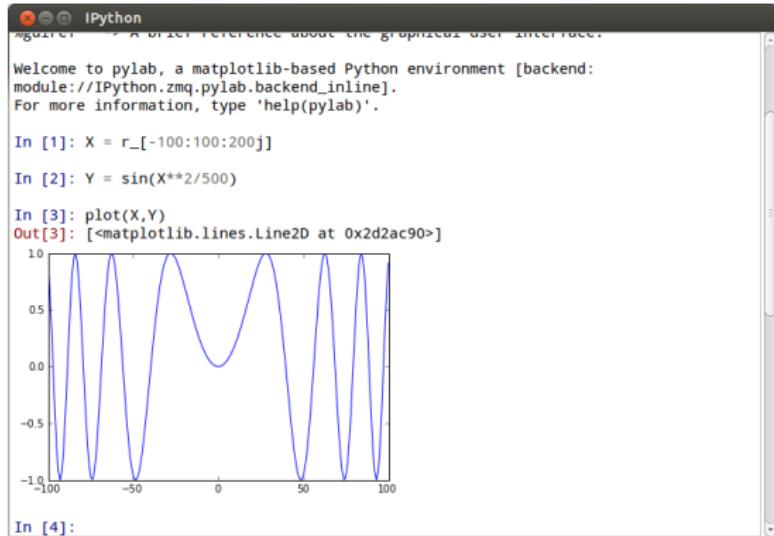
```
Python 2.7.3rc2 (default, Apr  6 2012, 21:13:11)
[GCC 4.6.3] on linux2
Type "copyright", "credits" or "license()" for more information
.
=====
>>> import numpy as np
>>> A = np.eye(5)
>>> print(A)
[[ 1.  0.  0.  0.  0.]
 [ 0.  1.  0.  0.  0.]
 [ 0.  0.  1.  0.  0.]
 [ 0.  0.  0.  1.  0.]
 [ 0.  0.  0.  0.  1.]]
>>> A[1:4,1:4] = np.ones((3,3))*5
>>> print(A)
[[ 1.  0.  0.  0.  0.]
 [ 0.  5.  5.  5.  0.]
 [ 0.  5.  5.  5.  0.]
 [ 0.  5.  5.  5.  0.]
 [ 0.  0.  0.  0.  1.]]
>>> |
```

Ln: 20 Col: 4

- The original IDE, bundled with most Python installations.
- IDLE is simple and not very intuitive to use.

More info: <http://docs.python.org/library/idle.html>

Ways to Use Python: IPython Qt Console



- IPython Qt Console allows for inline plotting.
- We will use IPython Qt Console to learn the basics of Python.

More info: <http://ipython.org/>

(Kitware / Google)

Python for MATLAB Users

June 2012

19 / 201

Ways to Use Python: SymPy with IPython Notebook

The screenshot shows the IPython Notebook interface running in Mozilla Firefox. The left sidebar contains navigation buttons for 'New', 'Open', 'Download', 'ipynb', 'Print', 'Cell', 'Kernel', and 'Help'. The main area displays the following code and its output:

```
In [16]: P = Matrix([[a, 0, px, 0],
   [0, a, py, 0],
   [0, 0, f/(f-n), -f*n/(f-n)],
   [0, 0, 1, 0]])
P
```

Out[16]:
$$\begin{pmatrix} a & 0 & p_x & 0 \\ 0 & a & p_y & 0 \\ 0 & 0 & \frac{f}{f-n} & -\frac{f_n}{f-n} \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Compute the Inverse of the Projection Matrix

```
In [17]: Pi = P.inv()
Pi
```

Out[17]:
$$\begin{pmatrix} \frac{1}{a} & 0 & 0 & -\frac{p_x}{a} \\
 0 & \frac{1}{a} & 0 & -\frac{p_y}{a} \\
 0 & 0 & \frac{1}{f-n} & \frac{n}{f-n} \\
 0 & 0 & -\frac{f-n}{fn} & \frac{1}{n} \end{pmatrix}$$

Map a world point (x,y,z) into normalized coordinates

- IPython Notebook lets you keep a math notebook in a web browser.
- SymPy adds symbolic math for a Maple or Mathematica alternative.
- Run `ipython notebook --profile=sympy`

More info: <http://ipython.org/> and <http://sympy.org/>



Ways to Use Python: Spyder IDE

The screenshot shows the Spyder IDE interface. The top menu bar includes File, Edit, Insert, Cell, Run, Help, and Spyder. The toolbar contains icons for file operations like Open, Save, and Print, along with other tools. The Project explorer on the left lists several Python files under 'Examples' and 'Exercises'. The Editor tab shows the code for 'HelloWorld.py':

```
1 #!/usr/bin/python
2 # Copyright 2011, Kitware, Inc.
3
4
5 """Prints the string "Hello World" and exits."""
6
7 my_text = "Hello World"
8
9 def main():
10     """The main function."""
11     print(my_text)
12     unused_var = 1
13
14
15 if __name__ == "__main__":
16     main()
17
18
19
```

The Variable explorer shows the current variables:

Name	Type	Size	Value
I	int32	(5, 5)	array([[1, 0, 0, 0, -1, 0, 1, 0, 0, 0], [0, 1, 0, 0, 0, 1, 0, 0, 0, 0], [-1, 0, 1, 0, 0, 0, 1, 0, 0, 0], [0, 0, 1, 0, 0, 0, 0, 1, 0, 0], [0, 0, 0, 1, 0, 0, 0, 0, 1, 0], [0, 0, 0, 0, 1, 0, 0, 0, 0, 1], [-1, 0, 0, 0, 0, 1, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 1, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 1, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0, 1, 0]])
X	Float64	(200,)	array([-100.0, -95.9 ...])
e	float	1	2.718281828459845
pi	float	1	3.141592653589793

The Object inspector panel shows information about the 'rand()' function. The Console tab displays the Python 1 interpreter output:

```
03:26:19
Type "help", "copyright", "credits" or "license" for more information.

Imported NumPy 1.6.1, SciPy 0.9.0, Matplotlib 1.1.1rc
Type "scientific" for more details.
>>> X = [-100:100:200j]
>>> I = eye(5)
>>> I = eye(5,dtype='int')
>>> rand(5,4)
array([ 0.32150613,  0.8368842 ,  0.83127323,  0.33114045],
      [ 0.25432276,  0.72797379,  0.26759263,  0.10712674],
      [ 0.23400013,  0.95567306,  0.67938578,  0.85042327],
      [ 0.53141453,  0.65109928,  0.57620572,  0.06570502],
      [ 0.06943613,  0.12310735,  0.25013823,  0.92263486])
>>>
```

At the bottom, the status bar shows permissions (RW), end-of-lines (LF), encoding (ASCII), line (17), and column (1).

- Spyder provides an IDE similar to MATLAB.
- We will use Spyder later in this tutorial.

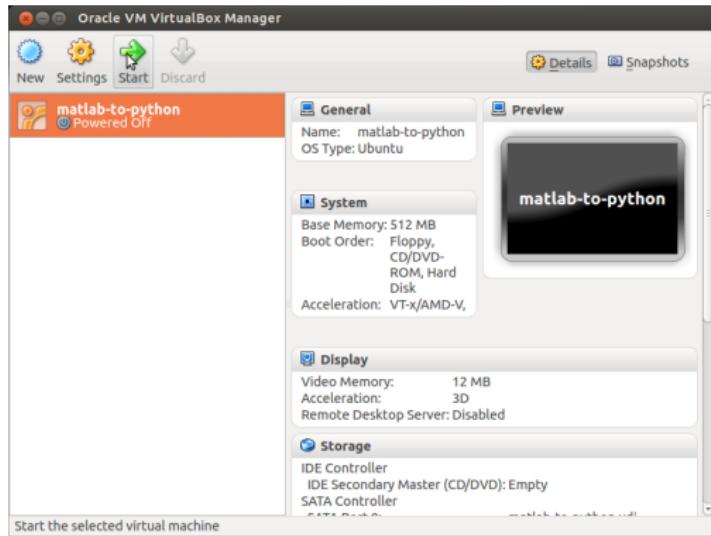
More info: <http://code.google.com/p/spyderlib/>

Using the Virtual Machine

Starting the Ubuntu 12.04 VM

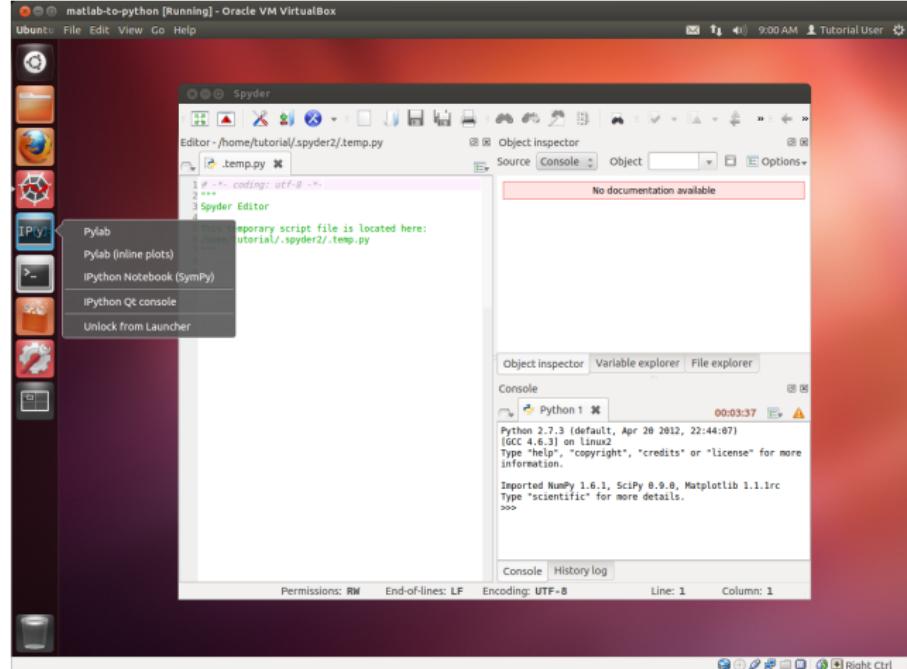
- Select the matlab-to-python VM.
- Press the “Start” button.
- Wait for the OS to boot.
- You will be automatically logged in as

username tutorial
password tutorial



The Ubuntu 12.04 Desktop

- The bar on the left is the App launcher.
- Some Apps have additional options when you right click.
- Drop down menus are on the top of the screen (as in MacOS X).



Key Applications

Key Applications



Firefox

A web browser.



Spyder

A MATLAB-like Python IDE.



IPython

An enhanced interactive Python shell.



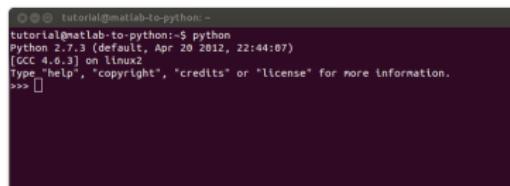
Terminal

The Gnome terminal emulator.

Introduction to Python

Playing with Python Interpreters

- IPython is a Python interpreter with additional interactive capabilities and more MATLAB-like features, discussed later.
- We will start with IPython to learn the basics and then move to Spyder when working with Python script files.
- Note: both Python and IPython can be run in a terminal with no need for a GUI.
 - To test this, click  in the launcher to open a terminal.
 - Type **python** to start Python and **ipython** to start IPython.
 - If your prompt is **>>>** you have a standard Python interpreter.
 - If your prompt is **In [1]:** you have an IPython interpreter.
 - Type **quit()** or **Ctrl-D** to exit, or close the terminal.



```
tutorial@matlab-to-python:~$ python
Python 2.7.3 (default, Apr 20 2012, 22:44:07)
[GCC 4.6.3] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> 
```

Python Interpreter



```
tutorial@matlab-to-python:~$ ipython
Python 2.7.3 (default, Apr 20 2012, 22:44:07)
Type "copyright", "credits" or "license" for more information.

IPython 0.12.1 -- An enhanced Interactive Python.
?          -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help       -> Python's own help system.
object?   -> Details about 'object', use 'object??' for extra details.

In [1]: 
```

IPython Interpreter

Launch the IPython Qt Console

- Open the IPython Qt Console by clicking  in the launcher.



The screenshot shows a window titled "IPython" with a dark gray header bar. The title bar has standard window controls (close, minimize, maximize) and the title "IPython". Below the title bar, the Python 2.7.3 startup message is displayed:
Python 2.7.3 (default, Apr 20 2012, 22:39:59)
Type "copyright", "credits" or "license" for more information.

IPython 0.12.1 -- An enhanced Interactive Python.
? -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help -> Python's own help system.
object? -> Details about 'object', use 'object??' for extra details.
%uiref -> A brief reference about the graphical user interface.

In [1]: |

Basic Math

- Try typing commands directly into the IPython interpreter.

```
In [1]: 2 + 2          # This is a comment
Out[1]: 4

In [2]: 2.0 + 2.0      # Floating point addition
Out[2]: 4.0

In [3]: 5.0 / 2        # Floating point division
Out[3]: 2.5

In [4]: 5 / 2          # Integer division (in Python3 this returns 2.5)
Out[4]: 2

In [5]: 5.0 // 2.0     # Double slash forces integer division
Out[5]: 2.0

In [6]: 5 % 2          # 5 modulo 2, in MATLAB this is mod(5,2)
Out[6]: 1

In [7]: 3 ** 2         # Exponentiation, in MATLAB this is 3^2
Out[7]: 9

In [8]: _ + 2          # Underscore is the last returned value, in MATLAB this is 'ans'
Out[8]: 11
```

More Examples: <http://docs.python.org/tutorial/introduction.html>

Basic Built-in Types

- Basic Numerical Types
 - `int` - Integer type with unlimited precision (in Python 3).
 - `float` - Floating point type (usually double precision).
 - `complex` - Complex number with real and imaginary float components.
 - Note: in Python 2 `int` has limited precision and separate type `long` has unlimited precision.
- Special Types
 - `bool` - Boolean with values `False` or `True`.
 - `None` - A null object, more about this later.
- Container Types
 - `str` - A string of characters.
 - `tuple` - A fixed sequence of objects.
 - `list` - A mutable sequence of objects.
 - `dict` - A dictionary mapping one set of objects to another.
 - Plus various lesser used containers.
- Additional numerical types and containers are provided by NumPy.

More info: <http://docs.python.org/library/stdtypes.html>

Type Conversions

- Types are easily converted using the syntax `type(object)`
- Type conversion even works to and from strings
- Some type conversions (e.g. complex to float) are not well defined and throw an exception.

```
In [9]: float(7)
Out[9]: 7.0

In [10]: int(3.14159)
Out[10]: 3

In [11]: str(28)
Out[11]: '28'

In [12]: float('5.5')
Out[12]: 5.5

In [13]: complex(4.5)
Out[13]: (4.5+0j)

In [14]: bool(0)
Out[14]: False

In [15]: float(True)
Out[15]: 1.0
```

More info: <http://docs.python.org/library/stdtypes.html>

Assignment

- Variables are created during assignment

```
In [16]: my_var  
-----  
Traceback (most recent call last):  
  File "<ipython console>", line 1, in <module>  
NameError: name 'my_var' is not defined  
  
In [17]: my_var = 10  
  
In [18]: my_var  
Out[18]: 10
```

- Other assignment examples

```
In [19]: x, y, z = 10, 20, 30      # multiple assignment, uses tuples  
  
In [20]: (y + z) / x  
Out[20]: 5  
  
In [21]: x += 5                  # equivalent to x = x + 5  
  
In [22]: y **= 2                  # equivalent to x = x * 2  
  
In [23]: x, y, z  
Out[23]: (15, 40, 30)
```

More Examples: <http://docs.python.org/tutorial/introduction.html>

Complex Numbers

- Python natively supports complex numbers.
 - Imaginary numbers are written with the `j` or `J` suffix.

```
In [24]: 1j * 1J
Out[24]: (-1+0j)

In [25]: complex(3, 2)          # can also create with the 'complex' function
Out[25]: (3+2j)

In [26]: a = (-1 + 3j) * 2j

In [27]: a
Out[27]: (-6-2j)

In [28]: a.real                # access the real part
Out[28]: -6.0

In [29]: a.imag                # access the imaginary part
Out[29]: -2.0

In [30]: abs(a)                # abs(a) == sqrt(a.real**2 + a.imag**2)
Out[30]: 6.324555320336759

In [31]: a.conjugate()         # the complex conjugate
Out[31]: (-6+2j)
```

More Examples: <http://docs.python.org/tutorial/introduction.html>

Boolean Operators

Python	Result
<code>x or y</code>	if x is false, then y, else x (short-circuit)
<code>x and y</code>	if x is false, then x, else y (short-circuit)
<code>not x</code>	if x is false, then <code>True</code> , else <code>False</code>

MATLAB	Result
<code>x y</code>	if <code>logical(x)</code> , then 1, else <code>logical(y)</code> (short-circuit)
<code>x && y</code>	if <code>logical(x)</code> , then <code>logical(y)</code> , else 0 (short-circuit)
<code>~x</code>	if <code>logical(x)</code> , then 0, else 1

• Examples

MATLAB

```
> 0.0 || 'hello'  
1  
> 10 && (5 || 6)  
1  
> ~ 5  
0
```

Python

```
> 0.0 or 'hello'  
'hello'  
> 10 and (5 or 6)  
5  
> not 5  
False
```

Comparisons

MATLAB	Python	Meaning
<	<	strictly less than
\leq	\leq	less than or equal
>	>	strictly greater than
\geq	\geq	greater than or equal
$=$	$=$	equal
$\sim=$	\neq	not equal
N/A	is	object identity
N/A	is not	negated object identity

- These comparisons always return a Boolean value.
- Comparison between different numeric types casts from `int` to `float` to `complex` as needed. For example `3.0 == 3` is `True`.
- Only “not equal” operators differ in syntax.
- The concept of object identity is not found in MATLAB.

More info: <http://docs.python.org/library/stdtypes.html>

Object Identity

- Every Python object has an identity. The function `id(object)` returns a unique identifying integer.
- The operation `x is y` is equivalent to `id(x) == id(y)`.
- Assignment allows multiple variables to refer to the same object.
- For immutable types (e.g. numeric types, strings, tuples) creation of new objects may or may not refer to same-valued exiting object.
- For mutable type (e.g. list, dict) creation always creates a new object.

```
In [32]: x, y = 3, 4.4

In [33]: x == y, x is y, id(x), id(y)
Out[33]: (False, False, 30294792, 32375168)

In [34]: x = 2 + 2.4

In [35]: x == y, x is y, id(x), id(y)      # here 'x is y' could be True or False
Out[35]: (True, False, 37311776, 32375168)

In [36]: x = y

In [37]: x == y, x is y, id(x), id(y)      # here 'x is y' must be True
Out[37]: (True, True, 32375168, 32375168)
```

More info: <http://docs.python.org/reference/datamodel.html>

Tuples

- Tuples are immutable sequences of objects of any type.
- Tuple items are comma separated and often enclosed in parentheses.
- A trailing comma is required to make a tuple of length one.

```
In [38]: t = 1, 'hello', (3.14, None), 2.5

In [39]: t[0]          # access item by index
Out[39]: 1

In [40]: t[0] = 2    # item assignment not allowed (immutable)
TypeError: 'tuple' object does not support item assignment

In [41]: t[1:3]       # access a slice of the objects (start:end)
Out[41]: ('hello', (3.14, None))

In [42]: len(t)       # the length of the tuple
Out[42]: 4

In [43]: 2.5 in t    # test for inclusion
Out[43]: True

In [44]: t + (5,)    # concatenate tuples
Out[44]: (1, 'hello', (3.14, None), 2.5, 5)

In [45]: (1, 2) * 3  # concatenate n shallow copies
Out[45]: (1, 2, 1, 2, 1, 2)
```

More info: <http://docs.python.org/library/stdtypes.html>

Strings

- Strings are immutable sequences of characters.
- Strings are enclosed in single or double quotes.
- The concatenation and indexing syntax is the same as Tuple.
- Strings have many built-in functions.

```
In [46]: s = "Hello World"

In [47]: s.upper()                      # to upper case
Out[47]: 'HELLO WORLD'

In [48]: s[3:5] + ' - ' * 3 + s[::-1]    # concatenation and slicing
Out[48]: 'lo- - - dlroW olleH'

In [49]: s.split()                      # string splitting
Out[49]: ['Hello', 'World']

In [50]: ' - '.join(('a', 'b', 'c'))    # join strings with a separator
Out[50]: 'a - b - c'

In [51]: s.replace("Hello", "Goodbye")   # substring replacement
Out[51]: 'Goodbye World'

In [52]: "item %05d has value %g" % (23, 3.14**2) # operator % works like sprintf
Out[52]: 'item 00023 has value 9.8596'
```

More info: <http://docs.python.org/library/string.html>

Lists

- Lists are mutable sequences of objects of any type.
- List items are comma separated and enclosed in square brackets.
- The concatenation and indexing syntax is the same as Tuple.

```
In [53]: v = [1, 'hello', (3.14, None), 2.5]

In [54]: v[0] = 2           # item assignment

In [55]: v.append(0.5)    # append items to the end

In [56]: del v[2]         # delete an item

In [57]: v
Out[57]: [2, 'hello', 2.5, 0.5]

In [58]: v.sort()         # sort items in place

In [59]: v
Out[59]: [0.5, 2, 2.5, 'hello']

In [60]: v = ([],)*3      # a tuple with 3 shallow copies of the same empty list

In [61]: v[0].append(1)   # append to the first list in the tuple

In [62]: v               # all lists update because they are shallow copies
Out[62]: ([1], [1], [1])
```

More info: <http://docs.python.org/library/stdtypes.html>

Dictionaries

- Dictionaries are mutable mapping of one object to another.
- Key:value pairs are comma separated and enclosed in curly braces.

```
In [63]: d = {'a' : 1, 'b' : 2, 'c' : 3}

In [64]: d['a']      # access a value by key
Out[64]: 1

In [65]: d['d'] = 4    # change a value or insert a new key:value

In [66]: d.keys()     # a list of dictionary keys
Out[66]: ['a', 'c', 'b', 'd']

In [67]: d.values()   # a list of dictionary values
Out[67]: [1, 3, 2, 4]

In [68]: d.items()    # a list of (key, value) tuples
Out[68]: [('a', 1), ('c', 3), ('b', 2), ('d', 4)]
```

IPython Commands

- IPython adds several interactive capabilities to standard Python.
 - Tab Completion: start typing and hit the Tab key to get a list of possible completions. Works for variables, functions, paths, etc.
 - History: access previous commands and their results
 - adds `__` and `___` for second and third to last results.
 - use the command `history` to list the history of commands.
 - use `In[n]` to access the `n`-th command string.
 - use `Out[n]` to access the `n`-th result.
 - Documentation: type `objectName?` to get info about `objectName`. Note, the standard Python `help(objectName)` also provides some documentation.
 - Auto Call: lets you write functions as statements.
For example, `abs x` calls `abs(x)`. MATLAB also allows this.
 - IPython provides many more features not covered here.

IPython Commands

- Examples of using IPython history

```
In [1]: a = 1
In [2]: a += 1
In [3]: history
1: a = 1
2: a += 1
3: %ip.magic("history ")
In [4]: a
Out[4]: 2
In [5]: In[2]
Out[5]: u'a += 1\n'
In [6]: exec In[2]      # like 'eval' in MATLAB, Python also has 'eval' for expressions
In [7]: a
Out[7]: 3
In [8]: Out[4] + Out[7]
Out[8]: 5
```

Python Scripting



Launch the Spyder IDE

- So far we've been interacting with an interpreter.
How do we write a script?



- Open the Spyder IDE by clicking in the launcher.
- Note: This version of Spyder (2.1.9) does not fully integrate with the installed version of IPython (0.12).
 - The standard Python shell in Spyder has been enhanced with some features of IPython, like tab completion.
 - IPython will run within Spyder, but certain capabilities, like plotting, do not work.
 - IPython integration is scheduled to be fixed in Spyder 2.2, planned for October 2012.
 - For now we will use the standard interpreter, the one with the >>> prompt.



Using the Spyder IDE

- Spyder has an editor, interpreter, variable explorer, object inspector

The screenshot shows the Spyder IDE interface with the following components:

- Project explorer:** Shows a tree view of files in the 'Tutorial' directory, including 'HelloWorld.py' which is selected.
- Editor:** Displays the code for 'HelloWorld.py'. The code prints "Hello World" and exits. It includes imports for sys, os, and random, and defines a main function that prints the text and exits.
- Object inspector:** Shows the documentation for the 'rand' function from the 'None' module. It describes generating random values in a given shape from a uniform distribution over [0, 1].
- Variable explorer:** Lists variables in the current namespace:

Name	Type	Size	Value
I	int32	(5, 5)	array([[1, 0, 0, 0, -[0, 1, 0, 0, -
X	Float64	(200,)	array([-100, -95.9 ...
e	Float	1	2.718281828459845
pi	Float	1	3.141592653589793
- Console:** Displays the Python 1 console output, showing the import of NumPy, SciPy, Matplotlib, and scientific mode, followed by a series of random numbers generated by the rand function.

Hello World

- The simple way

```
print "Hello World"
```

- Note: In Python 3 `print` is a function, not a statement.

```
print("Hello World") # Python 3 version
```

- A more complete example

Examples/HelloWorld.py

```
1 #!/usr/bin/python
2 # Copyright 2011, Kitware, Inc.
3
4
5 """Prints the string "Hello World" and exits."""
6
7 my_text = "Hello World"
8
9 def main():
10     """The main function."""
11     print my_text
12
13
14 if __name__ == "__main__":
15     main()
```

Hello World: line-by-line

- Comments. The first line is a shebang directive for Unix.

```
1 #!/usr/bin/python
2 # Copyright 2011, Kitware, Inc.
```

- Documentation string (docstring) for the script. Uses triple quotes.

```
5 """Prints the string "Hello World" and exits."""
```

- Assigns the “Hello World” string to a global variable.

```
7 my_text = "Hello World"
```

- A function definition with docstring. It calls the print function.

```
9 def main():
10     """The main function."""
11     print my_text
```

- Call the main function only if this file is run as a script (not imported).

```
14 if __name__ == "__main__":
15     main()
```

Running a Script

- In Spyder run the active script with Menu “Run→Run”, or press F5
- Adjust the configuration settings and press the “Run” button
- Future runs use the same configuration
- Change the run configuration with Menu “Run→Configure”, or press F6
- Alternatively, in a terminal window (from the home directory) run `python Tutorial/Examples/HelloWorld.py`



Conditionals

MATLAB

```
if a > 10
    a = 10;
end
```

Python

```
if a > 10:
    a = 10
```

- Syntax of basic conditionals is very similar.
- MATLAB uses keyword **end** to designate the end of a block.
- Python uses : to start a block and indentation to designate duration.
 - These rules apply generally to loops, functions, classes, etc.
 - **Indentation is important in Python!**
 - Choose any indentation, but consistent indentation defines a block.

MATLAB

```
if a > 10
    disp('a greater than 10')
elseif a < 5
    disp('a less than 5')
else
    disp('a is ', a)
end
```

Python

```
if a > 10:
    print 'a greater than 10'
elif a < 5:
    print 'a less than 5'
else:
    print 'a is', a
```

- Note, **elseif** in MATLAB becomes **elif** in Python.

Loops

MATLAB

```
x = 1;
while x < 10
    x = x * 2;
end
```

Python

```
x = 1
while x < 10:
    x = x * 2
```

- MATLAB and Python `while` loops are essentially the same.
- `break` and `continue` statements are the same in both languages.

MATLAB

```
X = [5 4 3 2 1];
for v = X
    disp(v)
end
```

Python

```
X = [5, 4, 3, 2, 1]
for v in X:
    print v
```

- MATLAB and Python `for` loops each iterate over a container.
- MATLAB uses symbol `=` while Python uses keyword `in`.

Generating Ranges

- MATLAB has special syntax to generate a range of numbers
 - `first:last`
 - `first:step:last`
- Python use the `range` function to achieve similar behavior
 - `range(end)`
 - `range(first, end)`
 - `range(first, end, step)`

MATLAB

```
> 0:4  
0 1 2 3 4  
> 2:5  
2 3 4 5  
> 6:-2:2  
6 4 2
```

Python

```
> range(5)  
[0, 1, 2, 3, 4]  
> range(2, 6)  
[2, 3, 4, 5]  
> range(6, 1, -2)  
[6, 4, 2]
```

- Note that `last` is inclusive while `end` is not. In MATLAB $\text{first} \leq i \leq \text{last}$, but in Python $\text{first} \leq i < \text{end}$.
- Python's `range` function is limited to integer types. NumPy adds the `arange` function which generalizes to floats.

Defining Functions

MATLAB

```
function [s, d] = sum_diff(x, y)
% SUM_DIFF Sum and diff of two values.
% x is the first value
% y is the second value

s = x + y;
d = x - y;
```

Python

```
def sum_diff(x, y):
    """Sum and diff of two values.

    x is the first value
    y is the second value
    """
    return x + y, x - y
```

- MATLAB handles return values by explicitly naming them.
 - returns zero, one, or multiple named variables.
- Python always returns exactly one object with the `return` statement.
 - a standalone or missing `return` statement returns `None`.
 - returning multiple comma-separated objects (as above) returns a tuple.

MATLAB

```
> [v1, v2] = sum_diff(7, 4)
v1 = 11
v2 = 3
> v3 = sum_diff(7, 4)
v3 = 11
```

Python

```
> v1, v2 = sum_diff(7, 4)
> v3 = sum_diff(7, 4)
> print "v1 = ", v1, ", v2 = ", v2, \
", v3 = ", v3
v1 = 11, v2 = 3, v3 = (11, 3)
```

Defining Functions – Optional Arguments

- MATLAB function arguments are all optional, but you can enforce the number of arguments and set default values by checking `nargin`.
- Python function arguments are required unless a default value is specified.
- All optional arguments must come after all required arguments.

MATLAB

```
function val = my_func(req, opt1, opt2)
if nargin < 3
    opt2 = 'default string';
elseif nargin < 2
    opt1 = 3.14;
end
disp(opt2)
val = req * opt1
```

Python

```
def my_func(req, opt1=3.14,
            opt2='default string'):
    print opt2
    return req * opt1
```

- Python optional arguments can be omitted and/or specified as `key=value` pairs in any order, e.g in the above example
`my_func(1, 3.14, 'hi')` is the same as `my_func(1, opt2='hi')`

Managing Libraries of Functions

- MATLAB and Python have different strategies for organizing and calling libraries of functions.
- In MATLAB, importing functions is implicit.
 - Each function must reside in a file of the same name.
 - When a function is called, the MATLAB path is searched to find the first matching function.
 - Dealing with functions in different packages with the same name is problematic.
- In Python, importing functions is explicit.
 - Any number of functions can reside in a file.
 - Modules, packages, and function must be explicitly imported when needed.
 - Name spaces and renaming allows for easy resolution of name conflicts.

Importing Packages, Modules, Functions, etc.

- The Python **import** statement is used to explicitly import packages, modules, classes, functions, and variables.

`module` is a python file (`.py`) that may contain multiple functions, variables, and classes.

`package` is a collection of modules in a directory. The directory must contain a file named `__init__.py`. Packages can be nested.

- The **import** statement initializes a module or package by executing the code in the module or `__init__.py`.
- Python searches current directory and Python path to find matching modules and packages.
- Each module or package is imported **only once** in a Python session, call `reload()` to force reloading.

An Example Python Package

Directory structure of a Python package

vision/	Top-level package
__init__.py	Initialize vision package
learning/	Subpackage for learning
__init__.py	
adaboost.py	
svm.py	
tracking/	Subpackage for tracking
__init__.py	
kalman.py	
features/	Subpackage for features
__init__.py	
sift.py	
harris.py	
canny.py	

Importing Modules and Packages

- If the `vision` directory is in the Python path then the following works.

```
import vision.tracking.kalman          # executes vision/tracking/kalman.py
vision.tracking.kalman.predict()        # call a function in kalman.py
```

- Generally `__init__.py` includes statements to import all the sub-packages or modules recursively. So the following usually works, but it imports the whole `vision` package.

```
import vision                      # executes vision/__init__.py
vision.features.harris.detect()    # call a function in harris.py
vision.tracking.kalman.predict()  # call a function in kalman.py
```

- The `as` keyword allows renaming during import.

```
import vision.tracking.kalman as kf
kf.predict()
```

Importing Specific Functions or Variables

- The `from` keyword allows importing into the current name space.

```
from vision.tracking.kalman import predict, correct
predict()
correct()
```

- Rename specific functions with `as`.

```
from vision.tracking.kalman import predict as my_pred, correct as my_corr
my_pred()
my_corr()
```

- Using `*` imports of everything into the current name space.

```
from vision.tracking.kalman import *
predict()
correct()
```

The Python Path

- The Python path is set with environment variable PYTHONPATH.
- Modify `sys.path` to make temporary changes during execution.

```
In [9]: import sys  
  
In [10]: sys.path.append("/home/tutorial")  
  
In [11]: sys.path  
Out[11]:  
['',  
 '/usr/bin',  
 '/usr/lib/python2.7',  
 ...  
 '/usr/lib/python2.7/dist-packages/IPython/extensions',  
 '/home/tutorial']
```

Exercise 1

Basic Python

- Open the file Exercises/ex1.py
- Write the missing function `sqrt_dict` to compute a dictionary mapping the first `n` integers to the square roots of their values.
- Hint: use the `sqrt` function from the `math` module.
- This example uses `sys.argv` to access a list of command line arguments.

Exercises/ex1.py

```
19      # parse the command line argument
20      # require exactly 1 argument (argv[0] is the script name)
21      if len(sys.argv) != 2:
22          print "usage: %s number" % sys.argv[0]
23          return
24
25      # convert the command line argument to an integer
26      num = int(sys.argv[1])
27      # call the function
28      values = sqrt_dict(num)
```

Exercise 1: Answer

Basic Python

- Here is one possible answer

Exercises/Answers/ex1-answer.py

```
6 import sys
7 import math
8
9
10 def sqrt_dict(n):
11     """Compute a dictionary mapping the first n integers to square roots."""
12     d = dict()
13     for i in range(n):
14         d[i] = math.sqrt(i)
15     return d
```

- An alternate one-line solution uses list comprehensions

Exercises/Answers/ex1-answer.py

```
22     return dict([(i, math.sqrt(i)) for i in range(n)])
```

Scientific Computing in Python

NumPy, SciPy, and Matplotlib

- NumPy adds basic MATLAB-like capability to Python:
 - multidimensional arrays with homogeneous data types
 - specific numeric data types (e.g. `int8`, `uint32`, `float64`)
 - array manipulation functions (e.g. `reshape`, `transpose`, `concatenate`)
 - array generation (e.g. `ones`, `zeros`, `eye`, `random`)
 - element-wise math operations (e.g. `add`, `multiply`, `max`, `sin`)
 - matrix math operations (e.g. inner/outer product, rank, trace)
 - linear algebra (e.g. `inv`, `pinv`, `svd`, `eig`, `det`, `qr`)
- SciPy builds on NumPy (much like MATLAB toolboxes) adding:
 - multidimensional image processing
 - non-linear solvers, optimization, root finding
 - signal processing, fast Fourier transforms
 - numerical integration, interpolation, statistical functions
 - sparse matrices, sparse solvers
 - clustering algorithms, distance metrics, spatial data structures
 - file IO (including to MATLAB .mat files)
- Matplotlib adds MATLAB-like plotting capability on top of NumPy.

Interactive Scientific Python (aka PyLab)

- PyLab is a meta-package that includes NumPy, SciPy, and Matplotlib.
- The easiest (and most MATLAB-like) way to work with scientific Python is to import everything into the global name space.

```
from pylab import *
```

- IPython can be started with the --pylab command line option.
- Spyder starts its standard Python interpreter with these scientific libraries imported, much like PyLab, but not PyLab *per se*.

Scripting with Scientific Python

- When writing scripts it is recommended that you:
 - only import what you need, for efficiency
 - import packages into namespaces, to avoid name clashes
- The community has adopted abbreviated naming conventions:

```
import numpy as np
import scipy as sp
import matplotlib as mpl
import matplotlib.pyplot as plt
```

- Some different ways of working with NumPy are

```
from pylab import *          # Import everything. Good for interactive work,
                           # but not recommended for writing scripts.

from numpy import eye, array  # Import only what you need into the global namespace.
from numpy.linalg import svd  # Verbose syntax for import statements,
                           # but simple syntax for calling functions.

import numpy as np            # Import a superset of what you need into a namespace.
np.linalg.svd(np.eye(3))      # More verbose syntax for calling functions.
```



NumPy

- Numpy provides a good website for mapping MATLAB syntax into equivalent NumPy syntax:
http://www.scipy.org/NumPy_for_Matlab_Users
- This tutorial will only cover the most common cases.

More info: http://www.scipy.org/NumPy_for_Matlab_Users



NumPy Array

- Use `numpy.array` to represent multidimensional arrays.

```
>>> A = array([1, 2, 3])      # create a one dimensional array from a list
>>> B = array([[1, 2, 3],
              [4, 5, 6]])  # create a two dimensional array from a list of lists
>>> C = array([B, B, B, B])  # create a 3D array from a list of 2D arrays
>>> A.shape                # the size of each dimension
(3,)
>>> B.shape
(2,3)
>>> C.shape
(4,2,3)
>>> A.ndim                  # the number of dimensions
1
>>> A.dtype                 # the data type of each element
dtype('int64')
```



NumPy Array

- Unlike MATLAB, operations on `numpy.array`s are element-wise.

MATLAB

```
> M = [1 2; 2 1];
> disp(M*M)
 5   4
 4   5

> disp(M.*M)
 1   4
 4   1
```

Python

```
> M = array([[1, 2], [2, 1]])
> print M*M
[[1 4]
 [4 1]]

> print dot(M,M)
[[5 4]
 [4 5]]
```

- The function `numpy.dot` is required for matrix (or tensor) multiplication.
- This syntax may seem ugly, but makes operations more consistent over arrays with dimensions greater than 2.



NumPy Matrix

- Numpy provides a special `numpy.matrix` subclass of `numpy.array`.
- `numpy.matrix` acts more like a MATLAB matrix.

MATLAB

```
> M = [1 2; 2 1];
> disp(M*M)
5     4
4     5
```

Python

```
> M = matrix('1 2; 2 1')
> print M*M
[[5 4]
 [4 5]]
```

- `numpy.matrix` can use a MATLAB-style string initializer.
- Use “`from numpy.matlib import *`” to override common array constructors (`zeros()`, `eye()`, etc.) with matrix versions.
- `numpy.matrix` only handles two dimensional arrays.
- `numpy.array` is more consistent and still preferred in most cases.



NumPy Numerical Data Types

- Numpy adds several numerical data types to Python

Names	Descriptions
<code>bool</code>	Boolean (<code>True</code> or <code>False</code>) stored in a byte
<code>int</code>	Platform integer (usually <code>int32</code> or <code>int64</code>)
<code>int8, int16, int32, int64</code>	8, 16, 32 and 64 bit signed integers
<code>uint8, uint16, uint32, uint64</code>	8, 16, 32 and 64 bit unsigned integers
<code>float</code>	Shorthand for <code>float64</code>
<code>float16, float32, float64</code>	half, single, and double precision floats
<code>complex</code>	Shorthand for <code>complex128</code>
<code>complex64, complex128</code>	complex numbers with single and double precision floats for real and imaginary parts

- Type can be specified during array creation and converted later

```
>>> A = np.array([1, 2, 3], dtype=np.float32)      # construct as type float32

>>> A
array([ 1.,  2.,  3.], dtype=float32)

>>> A.astype('complex')                          # convert the type of an existing array
array([ 1.+0.j,  2.+0.j,  3.+0.j])
```

More info: <http://docs.scipy.org/doc/numpy/user/basics.types.html>



NumPy Array Construction

- Common helper functions for constructing arrays are similar to MATLAB.

MATLAB

```
> M = zeros(3,4);
> M = ones(3,4);
> M = eye(3);
> M = diag([1 2 3]);
> M = diag(A);
> M = rand(3,4);
> M = repmat(A, m, n);
```

Python

```
> M = zeros((3,4))
> M = ones((3,4))
> M = eye(3)
> M = diag([1, 2, 3])
> M = diag(A)
> M = random.rand(3,4)
> M = tile(A, (m, n))
```



NumPy Array Slicing

- Slicing allow access to subset of arrays.
- Python slicing syntax is `first:end:step`.

MATLAB	NumPy	Description
<code>a(end)</code>	<code>a[-1]</code>	access last element in the 1xn matrix <code>a</code> .
<code>a(2,5)</code>	<code>a[1,4]</code>	access element in second row, fifth column.
<code>a(2,:)</code>	<code>a[1] or a[1,:]</code>	access entire second row of <code>a</code> .
<code>a(1:5,:)</code>	<code>a[0:5] or a[:5] or a[0:5,:]</code>	access the first five rows of <code>a</code> .
<code>a(3:2:21,:)</code>	<code>a[2:21:2,:]</code>	every other row of <code>a</code> , from the third to the twenty-first.
<code>a(1:2:end,:)</code>	<code>a[::2,:]</code>	every other row of <code>a</code> , starting with the first.
<code>a(end:-1:1,:)</code>	<code>a[::-1,:]</code>	<code>a</code> with rows in reverse order.



NumPy Indexing Tricks – Stacking

- Numpy provides a special class to enable several indexing tricks.
- The `r_` class provides row stacking of arrays
- `r_` is used with brackets as if it is being indexed.

```
>>> r_[[1, 2, 3], [4, 5, 6], [7, 8]]      # stack 1D arrays along the first axis
array([1, 2, 3, 4, 5, 6, 7])

>>> r_[eye(2), 2*ones((2,2)) ]           # stack 2D arrays along the first axis

array([[ 1.,  0.],
       [ 0.,  1.],
       [ 2.,  2.],
       [ 2.,  2.]])
```

- A variant named `c_` is same, but for column stacking.

```
>>> c_[eye(2), 2*ones((2,2)) ]           # stack 2D arrays along the second axis

array([[ 1.,  0.,  2.,  2.],
       [ 0.,  1.,  2.,  2.]])
```



NumPy Indexing Tricks – Range Generation

- The `r_` class also provides range generation.
- If an argument uses slice notation it is expanded to a range
- `r_[first:end:step]` is the same as `arange(first,end,step)`

```
>>> r_[5:10:0.5]      # generates a range from 5 to 10 with step 0.5
array([ 5. ,  5.5,  6. ,  6.5,  7. ,  7.5,  8. ,  8.5,  9. ,  9.5])
```

- When the step is replaced with an imaginary number the slice has a different interpretation, `first:last:num`

```
>>> r_[5:10:5j]      # generates 5 samples between 5 and 10, inclusive
array([ 5. ,  6.25,  7.5 ,  8.75, 10. ])
```



NumPy Indexing Tricks

- The `ix_` function generates indexes for irregular slices

MATLAB	NumPy	Description
<code>a([2,4,5],[1,3])</code>	<code>a[ix_([1,3,4],[0,2])]</code>	rows 2,4 and 5 and columns 1 and 3.

More info: http://www.scipy.org/NumPy_for_Matlab_Users

Matplotlib



- Matplotlib provides plotting capabilities that are similar to MATLAB.
- Syntax is the same or similar to MATLAB in most cases.
- Matplotlib has some advantages over MATLAB plots
 - Handles TeX equation typesetting within plots.
 - Easy export to PDF using vector graphics.
 - Plots often look nicer than MATLAB plots.
- The above logo is actually a matplotlib plot.

More info: <http://matplotlib.sourceforge.net/>

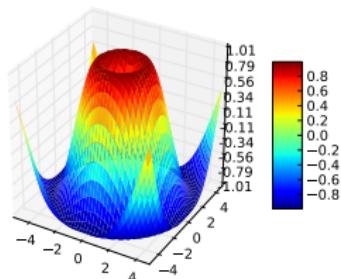
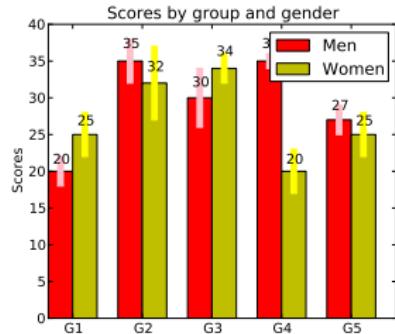
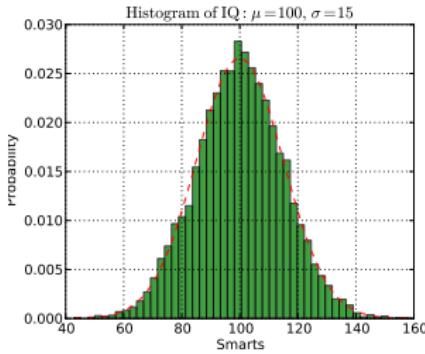
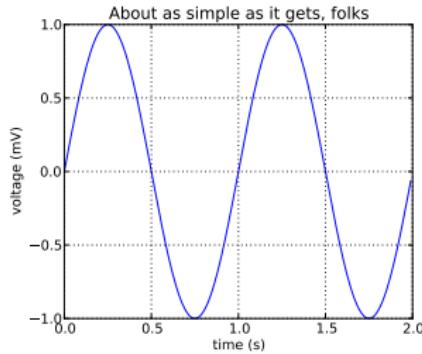
(Kitware / Google)

Python for MATLAB Users

June 2012

76 / 201

Matplotlib Example Plots

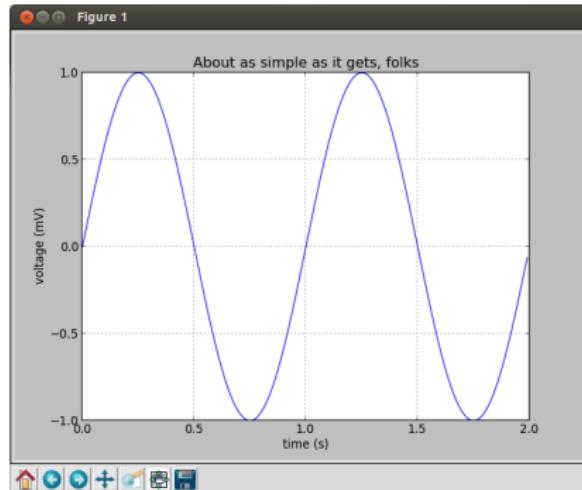


More info: <http://matplotlib.sourceforge.net/users/screenshots.html>

Matplotlib Example Plots

Examples/simple_plot.py

```
1 from pylab import *
2
3 t = arange(0.0, 2.0, 0.01)
4 s = sin(2 * pi * t)
5 plot(t, s, linewidth=1.0)
6
7 xlabel('time (s)')
8 ylabel('voltage (mV)')
9 title('About as simple as it gets, folks')
10 grid(True)
11 show()
```



- When not using pylab, these functions are found in `matplotlib.pyplot`.
- The figure is not displayed until `show()` is called.

More info: <http://matplotlib.sourceforge.net/users/screenshots.html>

Matplotlib Interactive Plotting

- Matplotlib has an interactive plotting mode that updates plots after each command, and `show()` is not needed.
- Interactive plotting is enabled with `ion()` and disabled with `ioff()`.
- Interactive mode is enabled automatically in Spyder and
`ipython --pylab`

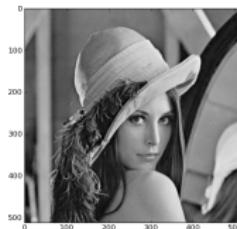
More info: <http://matplotlib.sourceforge.net/users/shell.html>

Exercise 2

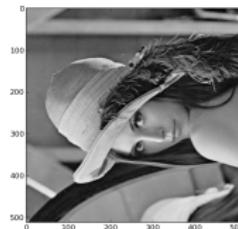
Plotting Lena with NumPy, SciPy, and Matplotlib

- ➊ Plot the “Lena” image.
- ➋ Plot the transpose of Lena using `lena().T`.
- ➌ Plot 100 random points on top of Lena with x and y between 0 and 512 using `rand` and `scatter`.
- ➍ Plot a histogram of Lena with 100 bins using `hist`. Hint: Flatten the image to 1-D using `lena().flatten()`
- ➎ Plot Lena with a sub-matrix from (200,200) to (400,400) flipped horizontally. Hint: Create a modified image array using slice notation.

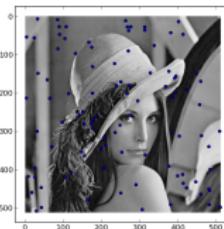
```
>>> from scipy.misc import lena  
>>> imshow(lena())  
>>> gray()
```



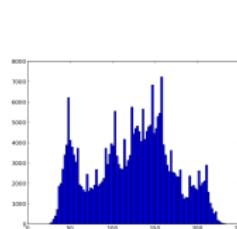
Part 1



Part 2



Part 3



Part 4



Part 5

Exercise 2: Answers

Plotting Lena with NumPy, SciPy, and Matplotlib

Exercises/Answers/ex2-answer.py

```
7  from scipy.misc import lena
8
9  figure("Part 1")
10 imshow(lena())
11 gray()
12
13 figure("Part 2")
14 imshow(lena().T)
15 gray()
16
17 figure("Part 3")
18 imshow(lena())
19 points = rand(2, 100) * 512
20 scatter(points[0], points[1])
21
22 figure("Part 4")
23 hist(lena().flatten(), 100)
24
25 figure("Part 5")
26 image = lena()
27 image[200:400, 200:400] = image[200:400, 400:200:-1]
28 imshow(image)
```



- Contains many scientific Python modules built on top of NumPy.
- Most interesting for computer vision work are
 - `ndimage` Multi-dimensional image processing
 - `linalg` Linear Algebra (extends numpy.linalg)
 - `fftpack` Fourier Transforms
 - `optimize` Optimization
 - `io` File IO (read/write MATLAB files)
- In this tutorial, we are only going to scratch the surface of what is contained in these libraries with a few examples.



- There are also many add-on toolkits for SciPy, collectively called SciKits
- Most interesting for computer vision work are
 - `scikits-image` Image Processing
 - `scikits-learn` Machine Learning and Data Mining
 - `ann` Approximate Nearest Neighbor library wrapper
 - `mlabwrap` Call MATLAB from Python



- Looking for an equivalent to MATLAB Central for sharing snippets of scientific Python code?
- Try SciPy Central (scipy-central.org)
- It's not as extensive as MATLAB Central, but it's fairly new.

More info: <http://scipy-central.org/>

(Kitware / Google)

Python for MATLAB Users

June 2012

84 / 201



- Let's work through an example of writing a Canny edge detector.
- This is not a very efficient implementation of Canny edge detection.
- The goal is to illustrate some SciPy and Numpy functionality.
- Open the file Examples/scipy_canny.py



- First import what we need: NumPy and and NDImage from SciPy.
- These modules are imported to abbreviated names `np` and `ndi`.
- Note that the constant `pi` (i.e. π) is imported into the global namespace to make math expressions easier to follow.

Examples/scipy_canny.py

```
4 import numpy as np
5 from numpy import pi
6 import scipy.ndimage as ndi
```



- Now look at the end of the file.

Examples/scipy_canny.py

```
46 if __name__ == "__main__":
47     from scipy.misc import lena
48     from matplotlib.pyplot import imshow, gray, show
49     imshow(canny_edges(lena()))
50     gray()
51     show()
```

- Usually all imports happen at the top of a file, but this is not required.
- In this example some more things are imported only when this file is run as a script.
- When run as a script, this code calls `canny_edges` on the Lena image and displays the result in a matplotlib GUI.
- When imported into another module, there is no dependency on matplotlib or Lena.

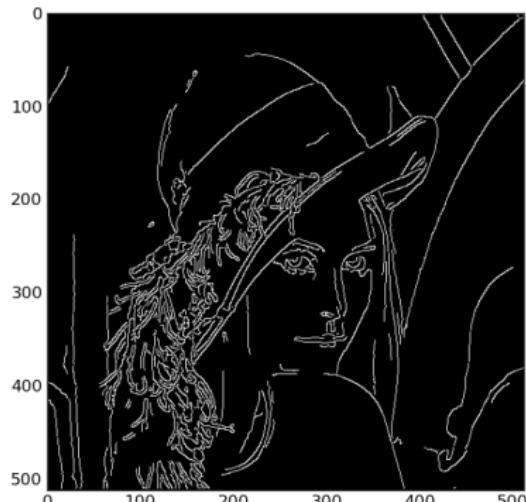
SciPy By Example – Result



- Calling `imshow(lena())` would show the input image on the left.
- Running this script displays the image on the right.



Input Image



Canny Edges



- Here is the Canny edge detection code.

Examples/scipy_canny.py

```
27 def canny_edges(image, sigma=1.0, low_thresh=50, high_thresh=100):
28     """Compute Canny edge detection on an image."""
29     image = ndi.filters.gaussian_filter(image, sigma)
30     dx = ndi.filters.sobel(image, 0)
31     dy = ndi.filters.sobel(image, 1)
32
33     mag = np.sqrt(dx ** 2 + dy ** 2)
34     ort = np.arctan2(dy, dx)
35
36     edge_map = non_maximal_edge_suppresion(mag, ort)
37     edge_map = np.logical_and(edge_map, mag > low_thresh)
38
39     labels, num_labels = ndi.measurements.label(edge_map, np.ones((3, 3)))
40     for i in range(num_labels):
41         if max(mag[labels == i]) < high_thresh:
42             edge_map[labels == i] = False
43     return edge_map
```

- Only `non_maximal_edge_suppression` is not provided by `np` or `ndi`.

SciPy By Example – Gaussian Filter

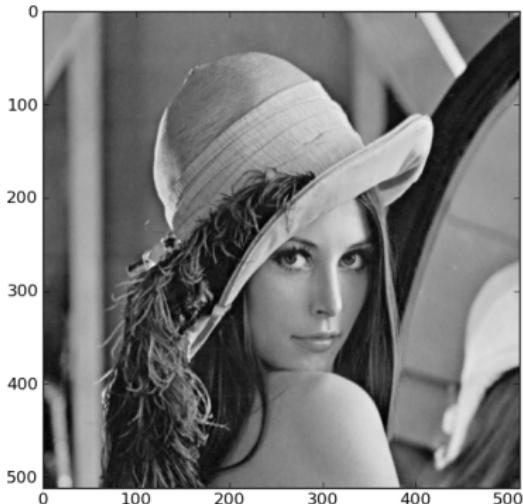


- Smooth the image with a Gaussian Filter.

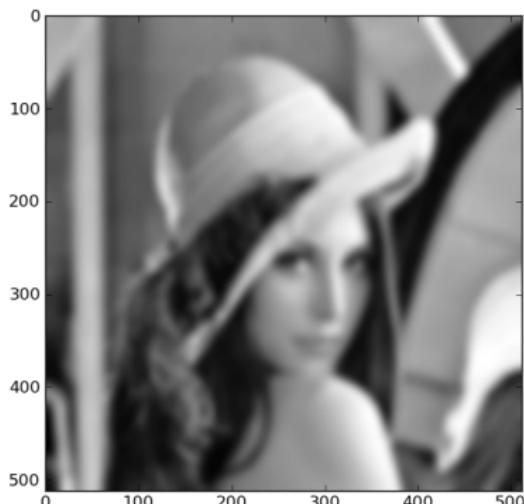
Examples/scipy_canny.py

29

```
image = ndi.filters.gaussian_filter(image, sigma)
```



Input Image



Gaussian Filter with $\sigma = 5$

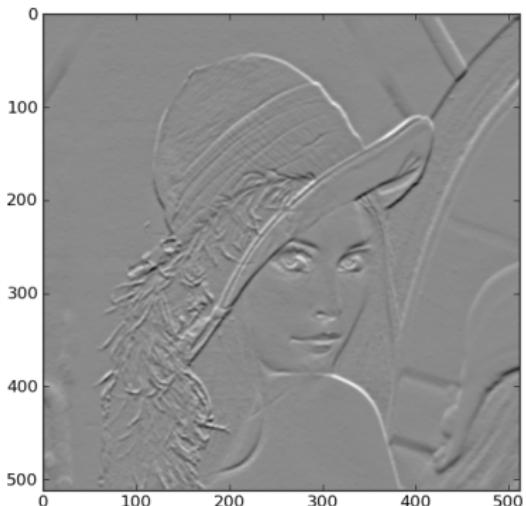
SciPy By Example – Sobel



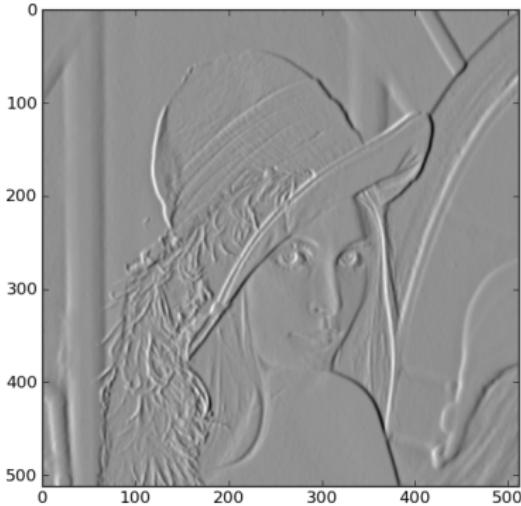
- ② Sobel derivatives on smoothed image.

Examples/scipy_canny.py

```
30     dx = ndi.filters.sobel(image, 0)
31     dy = ndi.filters.sobel(image, 1)
```



dx – Sobel horizontal



dy – Sobel vertical

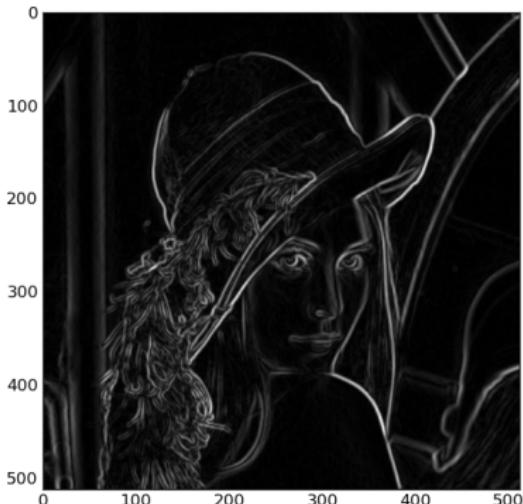


③ Sobel derivatives on smoothed image.

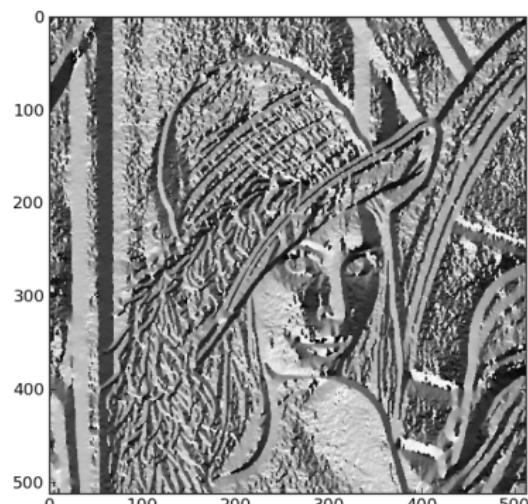
Examples/scipy_canny.py

33
34

```
mag = np.sqrt(dx ** 2 + dy ** 2)
ort = np.arctan2(dy, dx)
```



Gradient Magnitude



Gradient Orientation

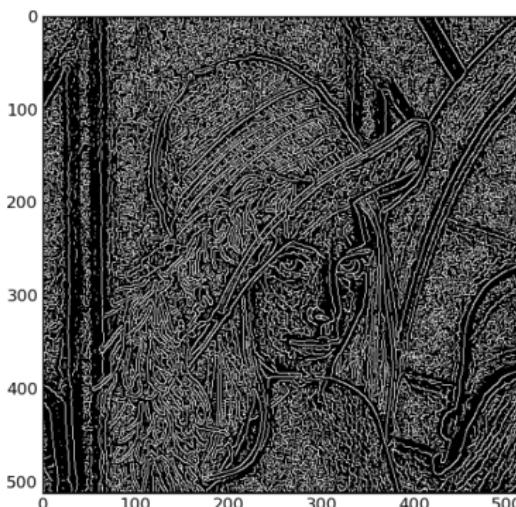


- ④ Non-maximal suppression (more on this later).

Examples/scipy_canny.py

36

```
edge_map = non_maximal_edge_suppresion(mag, ort)
```



All maximimal edges (binary image)

SciPy By Example – Threshold

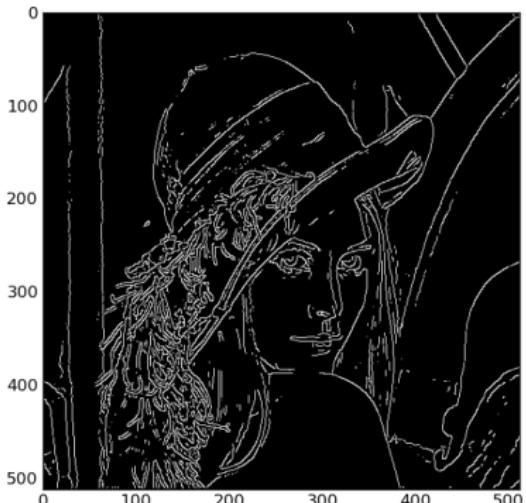


⑤ Apply the lower threshold

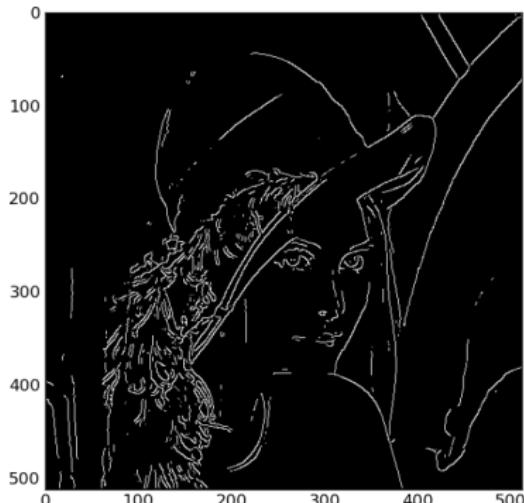
Examples/scipy_canny.py

37

```
edge_map = np.logical_and(edge_map, mag > low_thresh)
```



Low Threshold (50)



High Threshold (100)

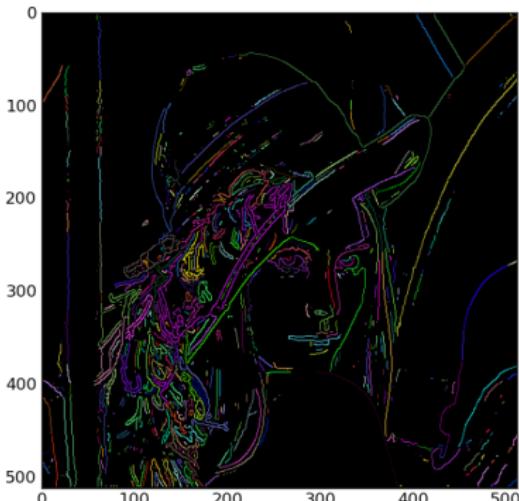


- ⑥ Label connected components using 8-neighborhood (3x3 mask of ones).

Examples/scipy_canny.py

39

```
labels, num_labels = ndi.measurements.label(edge_map, np.ones((3, 3)))
```



Connected edges with random colors

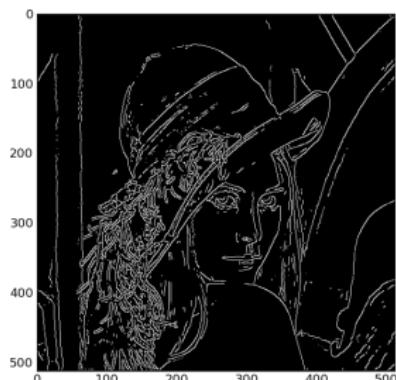
SciPy By Example – Hysteresis



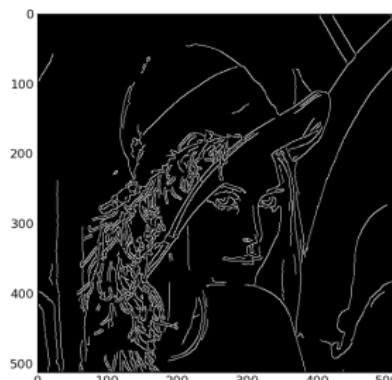
- 7 Remove connected edges not containing a value above the high threshold.

Examples/scipy_canny.py

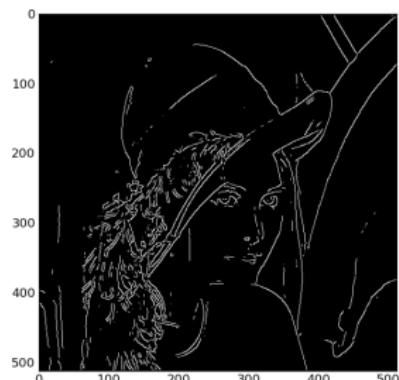
```
40     for i in range(num_labels):
41         if max(mag[labels == i]) < high_thresh:
42             edge_map[labels == i] = False
```



Low Threshold (50)



Hysteresis



High Threshold (100)



- Non-Maximal Suppression of edges produces a binary edge map.
- Gradient magnitude is a maximum in the direction normal to an edge.

Examples/scipy_canny.py

```
9 def non_maximal_edge_suppresion(mag, orient):
10     """Non Maximal suppression of gradient magnitude and orientation."""
11     # bin orientations into 4 discrete directions
12     abin = ((orient + pi) * 4 / pi + 0.5).astype('int') % 4
13
14     mask = np.zeros(mag.shape, dtype='bool')
15     mask[1:-1, 1:-1] = True
16     edge_map = np.zeros(mag.shape, dtype='bool')
17     offsets = ((1, 0), (1, 1), (0, 1), (-1, 1))
18     for a, (di, dj) in zip(range(4), offsets):
19         cand_idx = np.nonzero(np.logical_and(abin == a, mask))
20         for i, j in zip(*cand_idx):
21             if mag[i, j] > mag[i + di, j + dj] and \
22                 mag[i, j] > mag[i - di, j - dj]:
23                 edge_map[i, j] = True
24
25     return edge_map
```

Computer Vision in Python

Python Modules for Computer Vision

- This tutorial will focus on directly using Python interfaces to large open source C++ projects:
 - OpenCV
 - ITK
- Several packages for computer vision are built on top of some combination of NumPy, SciPy, OpenCV, and various other modules

SimpleCV <http://simplecv.org/>

PyVision <http://sourceforge.net/projects/pyvision/>

PyCVF <http://pycvf.sourceforge.net/>

PyCam <http://code.google.com/p/pycam/>

NDVision <https://launchpad.net/ndvision>

Loading and Saving Images with PIL

- The Python Imaging Library (PIL) has many image manipulation functions.
- PIL is most useful for loading and saving images in a large variety of standard image file formats.
- PIL uses its own image class, but it is easily converted to/from NumPy arrays.

```
import Image
import numpy as np

def load_image(file_name):
    """Read an image file and return a NumPy array."""
    return np.array(Image.open(file_name))

def save_image(data, file_name):
    """Save a NumPy array in an image file."""
    return Image.fromarray(data).save(file_name)
```



What is OpenCV?

- C/C++ library with Python wrappers
- Open source
- Very active community
- Implementations of many computer vision algorithms



- Graphical User Interface(GUI) - very easy procedural interface
- Input/Output - load/write images and movies in a variety of formats
- Image Processing - filtering,edges,histograms, etc.
- Machine Learning - SVMs, Decision Trees,
- Features
- Object Detection
- and MORE! (<http://docs.opencv.org/>)



- Two versions of wrappers: `cv` and `cv2`
- `cv2` contains `cv`
- `cv2` is preferred because it natively uses `numpy`

Examples/opencv_rand.py

```
3 import cv2
4 import numpy as np
```



Examples/opencv_rand.py

```
1 #!/usr/bin/python
2
3 import cv2
4 import numpy as np
5
6
7 a = np.random.rand(256, 256)
8 cv2.imshow('Random Image', a)
9 cv2.waitKey()
10 cv2.destroyAllWindows()
```



- Run Examples/opencv_rand.py



- Generate an array of random numbers with NumPy.

```
7 a = np.random.rand(256, 256)
```

- Show the array as an image in an OpenCV GUI window.
The string is both the window identifier and the window title.

```
8 cv2.imshow('Random Image', a)
```

- Wait for the user to press a key.
Optionally this function can take a timeout in milliseconds.

```
9 cv2.waitKey()
```

- Close the window and free memory. This is not strictly required in this case because it happens automatically when the script terminates.

```
10 cv2.destroyAllWindows('Random Image')
```



Examples/opencv_hello.py

```
1 #!/usr/bin/python
2
3 import cv2
4 import sys
5
6 if len(sys.argv) < 2:
7     # with no arguments, use this default file
8     filename = "/home/tutorial/Tutorial/Data/photo.jpg"
9 elif len(sys.argv) == 2:
10    filename = sys.argv[1]
11 else:
12     sys.exit("Expecting a single image file argument")
13
14 image = cv2.imread(filename)
15 print image.shape
16 image_small = cv2.resize(image, (800, 600))
17 textColor = (0, 0, 255) # red
18 cv2.putText(image_small, "Hello World!!!", (200, 200),
19             cv2.FONT_HERSHEY_PLAIN, 3.0, textColor,
20             thickness=4, linetype=cv2.CV_AA)
21 cv2.imshow('Hello World GUI', image_small)
22 cv2.waitKey()
23 cv2.destroyAllWindows()
```

OpenCV By Example – Image Processing



- Run Examples/opencv_hello.py Data/photo.jpg



OpenCV By Example – Image Processing



- Use OpenCV to read the image file into a NumPy array

```
14 image = cv2.imread(filename)
```

- Resize the image to 800×600 using bilinear interpolation

```
16 image_small = cv2.resize(image, (800, 600))
```

- Define a tuple to represent a color in BGR color space

```
17 textColor = (0, 0, 255) # red
```

- Draw text on the image in at location (200, 200) at a scale of 3.0

```
18 cv2.putText(image_small, "Hello World!!!", (200, 200),
19             cv2.FONT_HERSHEY_PLAIN, 3.0, textColor,
20             thickness=4, linetype=cv2.CV_AA)
```



Examples/opencv_movie.py

```
1 #!/usr/bin/python
2
3 import cv2
4 import sys
5
6 if len(sys.argv) < 2:
7     # with no arguments, use this default file
8     filename = "/home/tutorial/Tutorial/Data/movie.mov"
9 elif len(sys.argv) == 2:
10    filename = sys.argv[1]
11 else:
12     sys.exit("Expecting a single video file argument")
13
14 movie = cv2.VideoCapture(filename)
15 while True:
16     valid_frame, image = movie.read()
17     if not valid_frame:
18         break
19     cv2.imshow('OpenCV Video Example', image)
20     cv2.waitKey(5)
21
22 cv2.destroyAllWindows()
```

- Run Examples/opencv_movie.py Data/movie.mov

OpenCV By Example – Movie Player



- Open the video file

```
14     movie = cv2.VideoCapture(filename)
```

- Start an infinite loop for reading and displaying video frames

```
15     while True:
```

- Read one frame if available, if not exit the loop

```
16         valid_frame, image = movie.read()
17         if not valid_frame:
18             break
```

- Display the image in a GUI window, and wait for 5 milliseconds

```
19         cv2.imshow('OpenCV Video Example', image)
20         cv2.waitKey(5)
```



Examples/opencv_features.py

```
1 #!/usr/bin/python
2
3 import cv2
4 import sys
5
6 if len(sys.argv) < 2:
7     # with no arguments, use this default file
8     filename = "/home/tutorial/Tutorial/Data/photo.jpg"
9 elif len(sys.argv) == 2:
10    filename = sys.argv[1]
11 else:
12     sys.exit("Expecting a single image file argument")
13
14 image = cv2.imread(filename)
15 image_small = cv2.resize(image, (image.shape[1] / 4, image.shape[0] / 4))
16 image_gray = cv2.cvtColor(image_small, cv2.cv.CV_RGB2GRAY)
17 features = cv2.goodFeaturesToTrack(image_gray, 50, .01, 50)
18 features = features.reshape((-1, 2))
19 for x, y in features:
20     cv2.circle(image_small, (x, y), 10, (0, 0, 255))
21 cv2.imshow('OpenCV Features Example', image_small)
22 cv2.waitKey()
23 cv2.destroyAllWindows()
```

OpenCV By Example – Feature Detection



- Run Examples/opencv_features.py Data/photo.jpg



OpenCV By Example – Feature Detection



- Convert the color space from RGB to Gray.
Feature detection requires a gray-scale image.

```
16 image_gray = cv2.cvtColor(image_small, cv2.cv.CV_RGB2GRAY)
```

- Run the “Good Features to Track” algorithm to find 50 points.
The last two arguments are a threshold and minimum distance between points.

```
17 features = cv2.goodFeaturesToTrack(image_gray, 50, .01, 50)
```

- The features are a 3D array of shape (50, 1, 2).
Reshape it to (50, 2). The -1 means infer from data length.

```
18 features = features.reshape((-1, 2))
```

- At each feature location (x, y) draw a circle of radius 10 in red.

```
19 for x, y in features:  
20     cv2.circle(image_small, (x, y), 10, (0, 0, 255))
```



Exercise 3

OpenCV

- Open the file Exercises/ex3.py.
- Modify the code to run `goodFeaturesToTrack` on the movie.
- Then modify it to run `Canny` edge detection instead.

Exercises/ex3.py

```
16 while True:
17     valid_frame, image = movie.read()
18     if not valid_frame:
19         break
20
21     # Insert code here to process the image before displaying it
22
23     cv2.imshow('Exercise 3: OpenCV', image)
24     cv2.waitKey(5)
```



Exercise 3: Answer

OpenCV

Exercises/Answers/ex3-answer1.py

```
19     image_gray = cv2.cvtColor(image, cv2.cv.CV_RGB2GRAY)
20     features = cv2.goodFeaturesToTrack(image_gray, 50, .01, 50)
21     if features is not None:
22         for x, y in features.reshape((-1, 2)):
23             cv2.circle(image, (x, y), 10, (0, 0, 255))
24     cv2.imshow('Exercise 3: OpenCV', image)
```

Exercises/Answers/ex3-answer2.py

```
19     image_gray = cv2.cvtColor(image, cv2.cv.CV_RGB2GRAY)
20     edges = cv2.Canny(image_gray, 100, 120)
21     cv2.imshow('Exercise 3: OpenCV', edges)
```

- Introduction to ITK
- Introduction to SimpleITK
- Getting SimpleITK
- Basic Operations
- Advanced Operations
- Building Bridges with Other Libraries

What is ITK ?

- C++ Library
- Open Source (Apache 2.0 License)
- Generic Programming / C++ Templates
- Image Processing
- Image Segmentation
- Image Registration
- ITK's main foci are segmentation and registration, especially in the context of 3D images.

Funding

- Funded (mostly) by the US National Library of Medicine
- With contributions from
 - National Institute of Dental and Craniofacial Research
 - National Science Foundation
 - National Eye Institute
 - National Institute of Neurological Disorders and Stroke
 - National Institute of Mental Health
 - National Institute on Deafness and Other Communication Disorders
 - National Cancer Institute

History

- Started in 2000
- Developed by Companies and Universities
 - GE Corporate Research
 - Insightful
 - Kitware
 - UNC Chapel Hill
 - University of Pennsylvania
 - University of Utah

History

- Recent contributions by a larger community
 - Harvard - Brigham and Women's Hospital
 - University of Iowa
 - Georgetown University
 - INRA - France
 - German Cancer Research Center
 - ...and many others ...

Project Profile

- Total lines of code: 1,886,674
- Active developers: 56 (past 12 months)
- Total developers: 146 (full history of the project)
- Top 2% largest open source teams in the world
- Estimated cost of development: \$34.5 M¹

¹Ohloh: <http://www.ohloh.net/p/itk>

Funding

- First 10 years of development = \$15M
- Refactoring of ITKv4 in 2011 = \$5M
- Yearly maintenance supported by NLM

Simple ITK

SimpleITK provides a template-free layer on top of ITK that is automatically wrapped for Python.

Installing Simple ITK

- Two main options are building from source with Python wrappings
- Or easy_install (easy_install is much easier).

```
bash ~/ easy_install SimpleITK # It's that easy (you may have to be root).
```

- Let's walk through an example of basic Simple ITK usage from Examples/sitk_intro.py.

Creating Images with Simple ITK

- It's very easy to create images using SimpleITK.
- The images will be two-dimensional or three-dimensional based on the constructor

Examples/sitk_intro.py

```
4 import SimpleITK as sitk
5
6 # Create Images
7 image = sitk.Image(256, 256, 256, sitk.sitkInt16)
8 twoD = sitk.Image(64, 64, sitk.sitkFloat32)
```

- `sitk` is the module
- `Image` is the constructor for the `Image` class
- Height, width, depth (omit depth for 2D images)
- Datatype (more on this later)

What just happened?

Examples/sitk_intro.py

```
11 print image.GetPixel(0, 0, 0)
12 image.SetPixel(0, 0, 0, 1)
13 print image.GetPixel(0, 0, 0)
```

- Get the voxel value at [0,0,0]?
- Hmm, I don't like it, so set to 1
- What is the value at [0,0,0] now?

What just happened?

Examples/sitk_intro.py

```
16 print image[0, 0, 0]
17 image[0, 0, 0] = 10
18 print image[0, 0, 0]
```

Without warning, we sprinkled syntactic sugar on you!

- `image[0,0,0]` is shorthand for `image.GetPixel(0,0,0)`
- `image[0,0,0] = 10` is shorthand for `image.SetPixel(0,0,0,10)`

SimpleITK Pixel Types

- Simple ITK supports a wide array of pixel types
- To get a human readable assessment of the pixel type:

Examples/sitk_intro.py

```
21 print image.GetPixelIDTypeAsString()
```

SimpleITK Image Summary

- Images are created using `SimpleITK.Image (w, h, d, Type)`
- Images can be 2- or 3-dimensional
- Images can describe themselves
- Images have simple pixel accessors

Filtering with SimpleITK

- Basic operations are very simple
- Let's look at an image with a Gaussian subtracted from it:

Examples/sitk_intro.py

22

```
smooth = sitk.SmoothingRecursiveGaussian(image, 2.0)
```

- The output of SmoothingRecursiveGaussian is of type float
- The input image is signed short
- Most SimpleITK filters with 2 inputs require the same type
- We need to cast in order to difference the input image with the smoothed image

Introducing Cast

- The Cast Filter works as you would expect.

Examples/sitk_intro.py

```
25 print "Before: ", smooth.GetPixelIDTypeAsString()
26 smooth = sitk.Cast(smooth, image.GetPixelIDValue())
27 print "After: ", smooth.GetPixelIDTypeAsString()
28 sitk.Show(sitk.Subtract(image, smooth), "DiffWithGaussian")
```

Sizes and Indices

- Extracting a region of interest from a SimpleITK image is easy with `sitk.Extract`.

Examples/sitk_intro.py

```
31 size = [64, 64, 0]
32 start = [64, 0, 0]
33 sitk.Show(sitk.Extract(image, size, start), "Extracted")
```

More Advanced Filtering

- Now that we can do basic filtering, let's look at more advanced filters
- How about Otsu Thresholding?

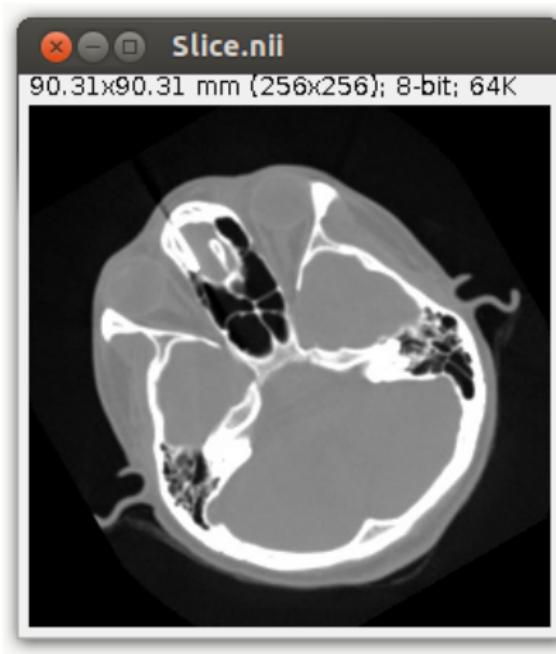
Examples/sitk_example.py

```
11 import SimpleITK as sitk
12
13 image = sitk.ReadImage(filename)
14 sitk.Show(image, 'Slice')
15 image = sitk.Cast(image, sitk.sitkFloat32)
16 thresh = sitk.OtsuThreshold(image, 0, 230, 15)
17 sitk.Show(thresh, 'Thresholded')
```

- Notice that filters can also be instantiated as objects

More Advanced Filtering

- We can run nearly every filter from ITK in this manner and quickly view results (as shown below).
- Bonus: `sitk.Show` allows us to view 3D Images with ImageJ



Integration with NumPy

- Let's look at integrating SimpleITK with NumPy.

Examples/sitk_numpy.py

```
11 import SimpleITK as sitk
12 import numpy as np
13 import pylab
14
15 image = sitk.ReadImage(filename)
16
17 a = sitk.GetArrayFromImage(image)
18
19 h, bins = np.histogram(a, bins=255, range=(0, 255))
20
21 mode = 0
22 for i in range(1, h.shape[0]):
23     if h[i] > h[mode]:
24         mode = i
25 print "Mode of image: " + str(mode)
26
27 # Print histogram
28 pylab.figure()
29 pylab.hist(a.flatten())
30 pylab.show()
```

Exercise 4

SimpleITK

- Open the file `Exercises/ex4.py`.
- Let's take the previous examples and build on them.
 - ① Use `sitk.RecursiveGaussian` to smooth the image, then plot the histogram with Matplotlib.
 - ② Use `sitk.Threshold` to segment bone from the input image.
- Hint: Bone has intensity greater than 210 in this image.

Exercise 4: Answer

SimpleITK

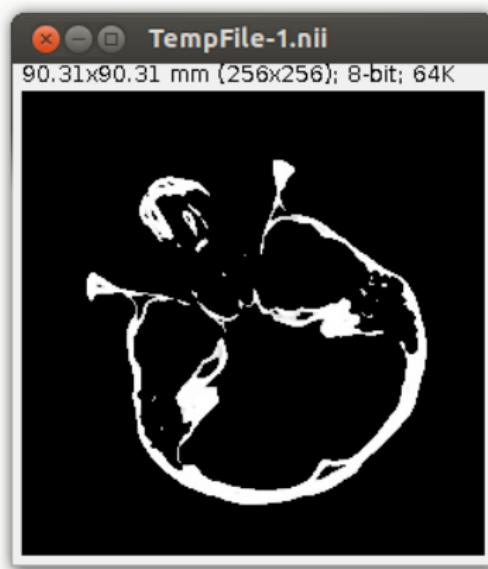
- Great Job! your code should look something like this

Exercises/Answers/ex4-answer.py

```
11 import SimpleITK as sitk
12 import pylab
13
14 image = sitk.ReadImage(filename)
15 image_array = sitk.GetArrayFromImage(image)
16
17 sitk.Show(image)
18
19 blurred = sitk.RecursiveGaussian(image, 4)
20 blurred_array = sitk.GetArrayFromImage(blurred)
21
22 # Generate a normal and blurred histogram
23 pylab.figure()
24 pylab.hist(image_array.flatten())
25 pylab.figure()
26 pylab.hist(blurred_array.flatten())
27 pylab.show()
28
29 # Threshold the image properly
30 thresh = sitk.Threshold(image, 210, 255)
31
32 sitk.Show(thresh)
```

More Advanced Filtering

- Your bone segmentation should look something like this.



Credits

Material taken from the following tutorials:

- The SimpleITK tutorial from MICCAI 2011 (<https://github.com/SimpleITK/SimpleITK-MICCAI-2011-Tutorial>).
- The ITKv4 next generation tutorial (<https://github.com/InsightSoftwareConsortium/ITKv4-TheNextGeneration-Tutorial>).

Parallel Processing in Python

Multi-Threading in Python

- Python supports multi-threading of code but has an important limitation.
- Only one thread can execute Python code at a time.
- Multi-thread can still be useful in certain cases such as
 - running multiple I/O-bound tasks simultaneously
 - code that uses a lot of C-based modules (e.g. NumPy and SciPy)
- To get parallel processing in Python it is often better to use multi-processing.

Multi-Processing in Python

- The `multiprocessing` module provides convenient ways to divide Python execution over multiple processes.
- Create a pool of worker processes and divide up the work.

Examples/multiproc.py

```
1  from multiprocessing import Pool
2
3  def f(x):
4      return x*x
5
6  if __name__ == '__main__':
7      pool = Pool(processes=4)          # start 4 worker processes
8      result = pool.apply_async(f, [10]) # evaluate "f(10)" asynchronously
9      print result.get(timeout=1)       # prints "100" unless your computer is slow
10     print pool.map(f, range(10))     # prints "[0, 1, 4,..., 81]"
```

Modern GPGPU

- GPUs have been evolving from dedicating graphics rendering hardware, to powerful general computing devices:

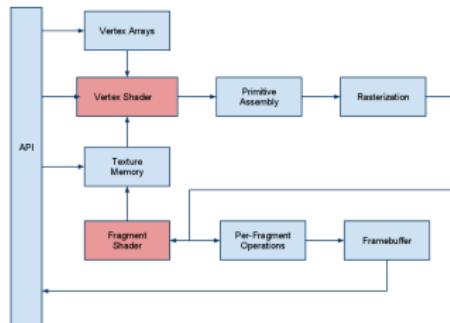


Modern GPU Architecture (Nvidia Kepler)

Modern GPGPU

- Massive amount of processing cores (3072 CUDA cores in Nvidia GeForce GTX 690 single card)
- Each core has its own dedicated register file
- Each cluster of cores share L1 cache
- Clusters of cores share L2 cache
- Huge memory throughput (512bit, 384GB/sec for Nvidia GeForce 690)
- Atomic operations
- 64-bit precision
- Error-correcting Code (ECC) memory to meet HPC requirement
- GPU Vs CPU on HPC performance
 - Most CV developers: 10 to 500x speed-up depending on the application and implementation
 - Intel: "GPUs are ONLY up to 14 times faster than CPU."

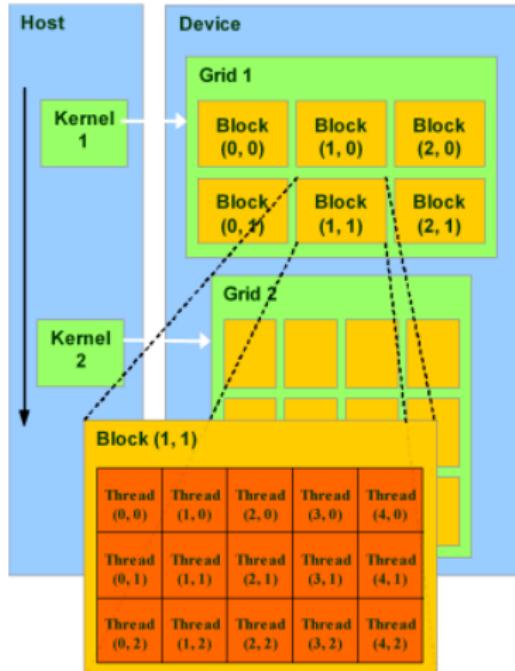
GPGPU Programming Model: Shading Language



Shading Language (Vertex & Fragment Shader)

- Map parallel computing task to OpenGL rendering pipeline
- Pack input data as vertices and texture, operate these data from vertex shader and fragment shader, results are output to frame buffer
- Pros:
 - enable the GPU for general use for the first time
- Cons:
 - single precision
 - no flow control and branching
 - very limited kernel size
 - lack of cross-kernel synchronization

GPGPU Programming Model: CUDA and OpenCL



- Enable C/C++ like direct parallel thread programming
- Enable cross-thread synchronization
- Enable direct memory access
- Enable atomic operations
- Effectively remove the kernel size limit
- Enable full support for flow control and branching
- GPUs become truly general purpose

CUDA Programming Model
(OpenCL has very similar model)

<http://ixbtlabs.com/articles3/video/cuda-1-p5.html>

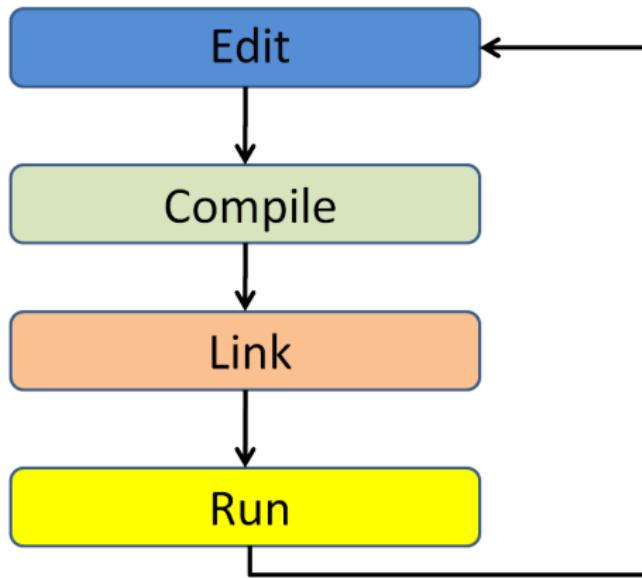
GPGPU in Computer Vision

- Many hardwired graphics rendering resources in GPU, including texture engine, 3D transformation, make GPU a natural fit for computer vision and computational photography processing.
- OpenVIDIA: Parallel GPU Computing Vision Library developed by Nvidia
- OpenCV: more and more GPU Modules
- Nvidia CUDA Zone: An Image and Computer Vision GPU library, including segmentation, feature processing, stereo imaging, machine learning and data processing

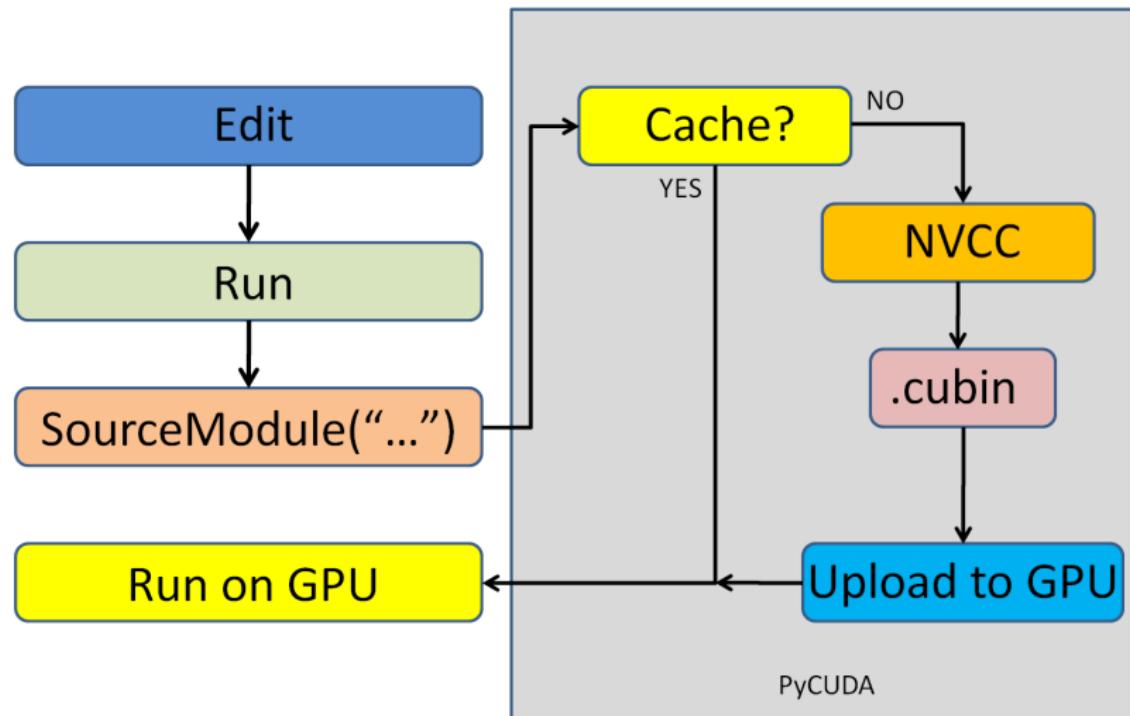
Why use GPGPU in Python?

- GPUs are everything that scripting languages are not
 - Highly parallel
 - Very architecture-sensitive
 - Built for maximum performance
- GPU and Python can be complement to each other
 - Use Python script as the brain, organizing processing module and data connection
 - Use GPU as muscles, executing intensive, dirty inner loops
- Play to the strengths of each programming environment
- Python + CUDA: PyCUDA
- Python + OpenCL: PyOpenCL

Native CUDA work flow



PyCUDA work flow



PyCUDA Example

```
import pycuda.driver as cuda          # import PyCUDA module
import pycuda.autoinit, pycuda.compiler    # initialize the CUDA hardware
import numpy

a = numpy.random.randn(4,4).astype(numpy.float32) # create a 4x4 array random numbers
a_gpu = cuda.mem_alloc(a.nbytes)                 # alloc the same sized array in GPU
cuda.memcpy_htod(a_gpu, a)                      # copy the data from host to device

mod = pycuda.compiler.SourceModule("""           # specify a CUDA kernel function
    __global__ void twice (float *a)
    {
        int idx = threadIdx.x + threadIdx.y*4;
        a[idx] *= 2;
    }
""")
func = mod.get_function("twice")                # get the handle of CUDA kernel
func(a_gpu, block=(4,4,1))                     # run the CUDA function with a configuration

a_doubled = numpy.empty_like(a)                  # create an empty buffer on host
cuda.memcpy_dtoh(a_doubled, a_gpu)              # copy result from device to host
print a_double
print a
```

gpuarray: Simple Linear Algebra

- Meant to look and feel just like numpy
- `+, -, *, /, fill, sin, exp, rand, basic indexing, norm, inner product,`
...
- Mixed types (`int32 + float32 = float64`)
- “`print gpuarray`” for debugging
- Allows access to raw bits

gpuarray: Element-wise expressions

- An example of element-wise operation in PyCUDA:

```
from pycuda.curandom import rand as curand      # import CUDA random number module

a_gpu = curand((50,))                           # create a 1-d array with random number
b_gpu = curand((50,))

from pycuda.elementwise import ElementwiseKernel # import ElementwiseKernel module

# specify the detail of element-wise operation
lin_comb = ElementwiseKernel(
    " float a, float *x, float b, float *y, float *z",
    "z[i] = a*x[i] + b*y[i]" )

c_gpu = gpuarray.empty_like(a_gpu)                # create a GPU array of same size
lin_comb(5, a_gpu, 6, b_gpu, c_gpu)              # run the ElementwiseKernel function
assert la.norm((c_gpu - (5*a_gpu+6*b_gpu)).get()) < 1e-5

print a_gpu
print b_gpu
print c_gpu
```

gpuarray: Reduction made easy

- Example: A scalar product calculation

```
from pycuda.reduction import ReductionKernel # import ReductionKernel module

# specify the detail of the reduction operation
dot = ReductionKernel(
    dtype_out=numpy.float32,
    neutral="0",
    reduce_expr="a+b",
    map_expr="x[i]*y[i]",
    arguments="const float *x, const float *y")

from pycuda.curandom import rand as curand

x = curand((1000*1000), dtype=numpy.float32)
y = curand((1000*1000), dtype=numpy.float32)

x_dot_y = dot(x, y).get()
x_dot_y_cpu = numpy.dot(x.get(), y.get())

print x
print y
print x_dot_y
print x_dot_y_cpu
```

PyCUDA: Vital Information

- <http://mathematician.de/software/pycuda>
- Complete documentation
- MIT License (no warranty, free for all use)
- Requires: numpy, Python 2.4+ (Win/OS X/Linux)
- Support via mailing list



<http://mathematician.de/software/pycuda>

(Kitware / Google)

Python for MATLAB Users

June 2012

156 / 201

Introducing... OpenCL and PyOpenCL

- OpenCL: the Open Computing Language
 - Vendor-independence
 - Single abstraction works well for GPUs, CPUs
 - A JIT C compiler baked into a library
 - Intel's future integrated GPUs will be supporting (more ubiquitous, available in cloud)
- PyOpenCL: the “PyCUDA” for OpenCL
 - Complete, mature API wrapper
 - Has: Arrays, element-wise operations, ...
 - <http://mathematician.de/software/pyopencl>



OpenCL: Same flavor, different recipe

```
import pyopencl as cl , numpy      # import PyOpenCL module

a = numpy.random.rand(50000).astype(numpy.float32)  # create a buffer on host

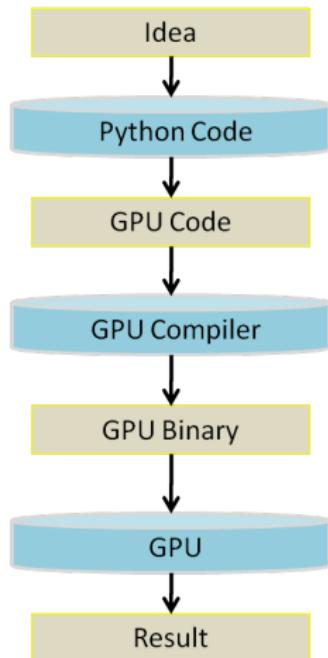
ctx = cl.create_some_context()      # create an OpenCL context
queue = cl.CommandQueue(ctx)

# create a buffer on device and copy the data in host buffer to device buffer
a_buf = cl.Buffer(ctx , cl.mem_flags.READ_WRITE, size=a.nbytes)
cl.enqueue_write_buffer(queue, a_buf , a)

# specify an OpenCL kernel function
prg = cl.Program(ctx, """
    __kernel void twice(__global float *a)
    {
        int gid = get_global_id (0);
        a[gid] *= 2;
    }
""").build()

# run an OpenCL kernel function
prg.twice(queue, a.shape, None, a_buf).wait()
```

Metaprogramming



- GPU code does not need to be a compile-time constant
- Code is data - it wants to be reasoned about at run time
- Automated tuning
- Data types
- Specialized code for given problem
- Constants faster than variables
- Loop unrolling

Showcase...

Interfacing Python with Other Languages

Wrapping C/C++ code in Python

There are several options for wrapping C/C++ code in Python

- Calling the Python C API Manually (not for the faint of heart)
- Interfacing with shared libraries using ctypes
- Using Cython to generate C API calls (good for optimizing critical pieces of code)
- Using SWIG to generate C API calls (good for wrapping larger libraries)
- Using Boost.Python to bridge between Python and C++

Working with the Python C API

- Calling the C API directly allows you to get the most performance possible and have the greatest control of execution.
- This is also one of the least easily maintainable methods.
- For more information, see
<http://docs.python.org/extending/extending.html>.
- There is a simple example at
Examples/c-api-demonstration/cdemo.c

A Note on the ctypes Module

- ctypes is a built-in module that allows calling functions from shared libraries.
- A more exhaustive explanation of what ctypes can do can be found here: <http://docs.python.org/library/ctypes.html>.
- For example, we can easily call the standard libc time function with the following snippet of code:

```
7  from ctypes import cdll, CDLL
8
9  # The name of our library
10 library_name = 'libc.so.6'
11
12 # Load the library as a module
13 cdll.LoadLibrary(library_name)
14 libc = CDLL(library_name)
15
16 # Call the standard time function and print the results
17 print libc.time(None)
```

- This can be very useful, but doesn't feel as "pythonic" as some of the other approaches.

Cython

- Cython is a tool for “compiling” Python-like code into C to be compiled into python extension modules.
- There are sets of annotations that can further optimize code through static typing, removal of bounds checking, etc.
- An example of a pyx file and generated C file can be found in Examples/cython
- Let’s go through a simple example.

Cython

- This simple code compiled with Cython is 35% faster.

```
0 def f(x):
1     return x**2-x
2
3 def integrate_f(a, b, N):
4     s = 0
5     dx = (b-a)/N
6     for i in range(N):
7         s += f(a+i*dx)
8     return s * dx
```

- By adding type declarations we can go a great deal faster, even with this simple example.

Cython

- This code compiled with Cython is 400% faster, but we lose compatibility with Python through the calls to cdef.

```
0 def f(double x):
1     return x**2-x
2
3 def integrate_f(double a, double b, int N):
4     cdef int i
5     cdef double s, dx
6     s = 0
7     dx = (b-a)/N
8     for i in range(N):
9         s += f(a+i*dx)
10    return s * dx
```

- We can do even better by typing the function call.

Cython

- This code compiled with Cython is 15,000% faster, but the calls are no able to raise Python errors.

```
0 cdef double f(double x) except? -2:
1     return x**2-x
2
3 def integrate_f(double a, double b, int N):
4     cdef int i
5     cdef double s, dx
6     s = 0
7     dx = (b-a)/N
8     for i in range(N):
9         s += f(a+i*dx)
10    return s * dx
```

Cython

- Let's take a look at that timing on our virtual machines using

```
In [1]: import simple
In [2]: import simplec
In [3]: import simple_opt
In [4]: import simple_opt_full
In [5]: %timeit for x in range(1,1001): simple.integrate_f(3, 5, x)
10 loops, best of 3: 137 ms per loop

In [6]: %timeit for x in range(1,1001): simplec.integrate_f(3, 5, x)
10 loops, best of 3: 79.8 ms per loop

In [7]: %timeit for x in range(1,1001): simple_opt.integrate_f(3, 5, x)
10 loops, best of 3: 55.6 ms per loop

In [8]: %timeit for x in range(1,1001): simple_opt_full.integrate_f(3, 5, x)
1000 loops, best of 3: 1.85 ms per loop
```

Cython

- Here is an example of how to compile your cython into a loadable module (taken from <http://docs.cython.org/src/reference/compilation.html>).

```
0 $ cython yourmod.pyx
1 $ gcc -shared -pthread -fPIC -fwrapv -O2 -Wall -fno-strict-aliasing \
2 -I/usr/include/python2.7 -o yourmod.so yourmod.c
```

- This should give you an idea of the possibilities for rapid optimization with Cython. See <http://cython.org/> for more information.

SWIG

- SWIG is a great option for wrapping existing C/C++ code for many different languages, including Python
- An example of swig usage can be found in `Examples/swig`
- SWIG is generally used for large projects that need semi-automated wrapping for multiple languages (SimpleITK uses this for its wrapping).

More Info: <http://www.swig.org/>

(Kitware / Google)

Python for MATLAB Users

June 2012

171 / 201

SWIG

- Here is an example of a C file we want to automatically wrap with SWIG

```
3 #include <time.h>
4 double My_variable = 3.0;
5
6 int fact(int n) {
7     if (n <= 1) return 1;
8     else return n*fact(n-1);
9 }
10
11 int my_mod(int x, int y) {
12     return (x%y);
13 }
14
15 char *get_time()
16 {
17     time_t ltime;
18     time(&ltime);
19     return ctime(&ltime);
20 }
```

SWIG

- In SWIG we define the interface or .i files

```
2 %module example
3 %
4 /* Put header files here or function declarations like below */
5 extern double My_variable;
6 extern int fact(int n);
7 extern int my_mod(int x, int y);
8 extern char *get_time();
9 %

10
11 extern double My_variable;
12 extern int fact(int n);
13 extern int my_mod(int x, int y);
14 extern char *get_time();
```

SWIG

- To generate our example wrapping and build example.c itself, we run the following commands.

```
3 swig -python example.i
4 gcc -c example.c example_wrap.c -I/usr/include/python2.7
5 ld -shared example.o example_wrap.o -o _example.so
```

- Swig generates a shared library and python wrapper for importing it as shown below.

```
In [9]: import example
In [10]: example.fact(10)
Out[10]: 3628800
```

- This pattern can be used to wrap existing C and C++ libraries in a maintainable way.
- There is also support for converting between STL types and Python types.

Boost.Python

- Boost.Python is an alternative to SWIG.
 - Only for making Python bindings (not multi-language)
 - Designed to make wrapping of C++ easy.
 - Handles C++ classes, virtual functions, inheritance, etc.
 - Only involves writing C++ code (no interface files)
- To wrap the previous example code

```
#include <boost/python.hpp>

BOOST_PYTHON_MODULE(example)
{
    using namespace boost::python;
    scope().attr("My_variable") = My_variable;
    def("fact", fact);
    def("my_mod", my_mod);
    def("get_time", get_time);
}
```

Interfacing with Matlab

Interfacing with Matlab

Any transition to Python from Matlab will require a means to utilize your previous Matlab code base. There are several methods for doing this:

① Converting Matlab Code to Python

- Open-Source Matlab-to-Python Compiler (OMPC) (ompc.juricap.com):
 - Converts existing Matlab functions to Python, does not require Matlab license, but only works for simple functions

② File Interchange with Matlab

- Load/Save .mats in Python:
 - Scipy.io.<*loadmat/savemat*> load/save .mats, does not require Matlab license, but may need to manipulate format after loading

③ Calling Legacy MATLAB Code from Python

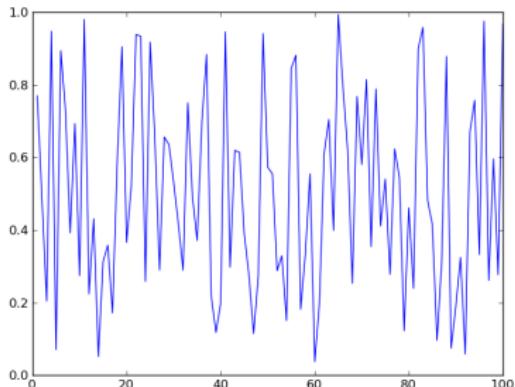
- Mlabwrap(mlabwrap.sourceforge.net):
 - Calls Matlab functions in manner similar to Matlab, does require Matlab license, and slower than Matlab

OMPC Example

Convert a Simple Matlab function that plots random numbers to Python

Convert .m to .pym

```
In [11]: import ompc  
In [12]: import plot_rand_nums_ex1  
In [13]: plot_rand_nums_ex1(100)
```



<http://ompc.juricap.com/download>

(Kitware / Google)

Python for MATLAB Users

June 2012

179 / 201

Examples/plot_rand_nums_ex1.m

```
1 function plot_rand_nums_ex1(maxX)  
2 %Function to plot a vector of random numbers  
3 x= 1:maxX;  
4 randx= rand(1,maxX);  
5  
6 plot(x,randx)
```

Examples/plot_rand_nums_ex1.pym

```
1 # This file was automatically translated  
2 # by OMPC (http://ompc.juricap.com)  
3  
4 from ompc import *  
5  
6 @mfunction("")  
7 def plot_rand_nums_ex1(maxX=None):  
8     #Function to plot a vector of random numbers  
9     x = mslice[1:maxX]  
10    randx = rand(1, maxX)  
11  
12    plot(x, randx)
```

Load/Save .mat Examples

Create Structure in Matlab, save as .mat, and load into Python

MATLAB

```
>> xx.range= 1:11;
>> xx.name= 'vars1';
>> xx.cellvalues= {[15 16], 'var1', [17 18], 'var2'}
xx =
    range: [1x11 double]
    name: 'vars1'
    cellvalues: {1x4 cell}
>> save('xx.mat','xx')
```

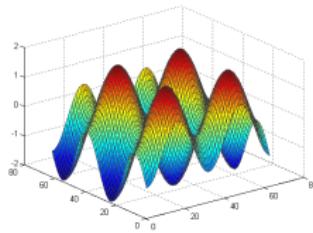
Python

```
> import scipy.io as sio
> xx= sio.loadmat('xx.mat')
> print xx
{'xx': array([[ [[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]], [u'vars1'],
[[array([[15, 16]]], dtype=uint8), array([u'var1'],
dtype=<U4>), array([[17, 18]]], dtype=uint8), array([u'var2'],
dtype=<U4>)]]], dtype=[('range', '|O4'), ('name', '|O4'), ('cellvalues', '|O4')]),
'__version__': '1.0', '__header__': 'MATLAB 5.0 MAT-file, Platform: PCWIN64,
Created on: Tue May 29 12:53:43 2012', '__globals__': []}}
```

Mlabwrap Example

- Call Matlab's Surf() function from Python

```
1  #!/usr/bin/python
2  # Copyright 2011, Kitware, Inc.
3  # http://mlabwrap.sourceforge.net/
4  ''' Examples for calling matlab functions from Python'''
5
6  from mlabwrap import mlab; from numpy import *
7
8  def main():
9      '''Calling Matlab's surf() command'''
10     xx = arange(-2*pi, 2*pi, 0.2) #Define range of inputs
11     mlab.surf(subtract.outer(sin(xx),cos(xx))) #call matlab's surf() function
```



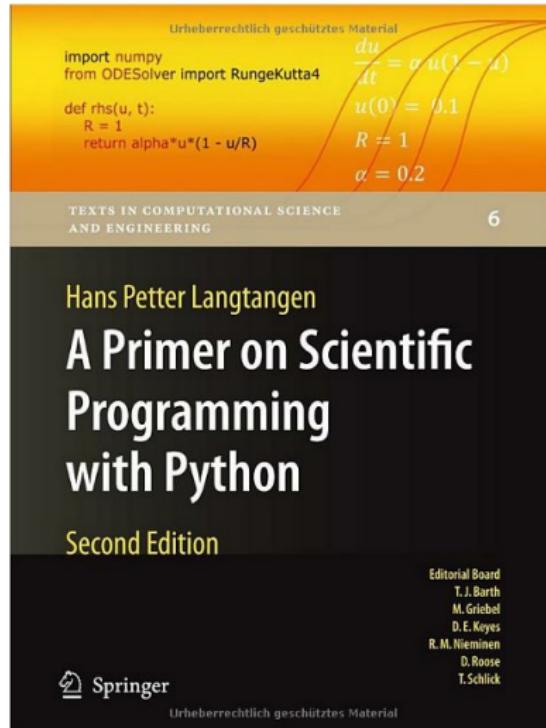
- Call Matlab's Singular Value Decomposition (svd) function

```
13  '''Calling Matlab's Singular Value Decomposition svd() command'''
14  U, S, V = mlab.svd([[1,2],[1,3]], nout=3)
```

More Info: <http://mlabwrap.sourceforge.net/>

Wrap-Up

Raffle Prize 1



Book: A Primer on Scientific Programming with Python

Raffle Prize 2



ASUS Transformer Pad TF300

Send Us Feedback

Please send us feedback, either positive or negative.

Thank You!



matt.leotta

amitha.perera patrick.reynolds eran.swears

@kitware.com

yongzhao

varun

@google.com

Appendix

Native Installation of Scientific Python

Native Installation of Scientific Python

- The virtual machine provides a nice way to try out software without worrying about installation.
- Eventually you'll want to install these toolkit in your native OS.
- The following slides provide basic installation instructions for
 - Windows
 - Mac OS X
 - Linux

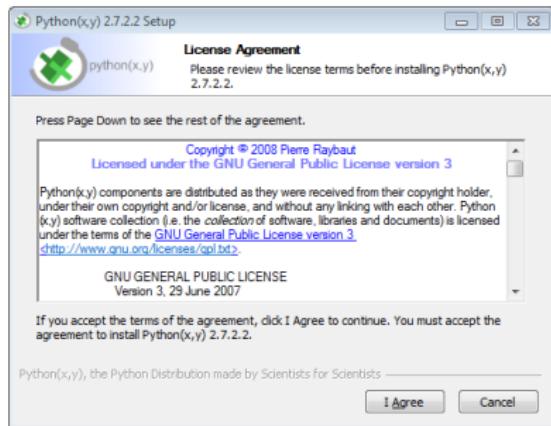
Installing on Windows

- It is difficult to install individual toolkits on Windows.
- Packages installed independently often do not work together because
 - package versions may be incompatible with each other
 - packages are compiled with different compilers.
- Luckily there is a project called Python(x,y) which installs almost everything you need for scientific Python and configures it to work together.



Installing on Windows – Python(x,y)

- ① Download the package installer (currently Python(x,y)-2.7.2.2.exe)
- ② Run the installer executable and accept the license agreement

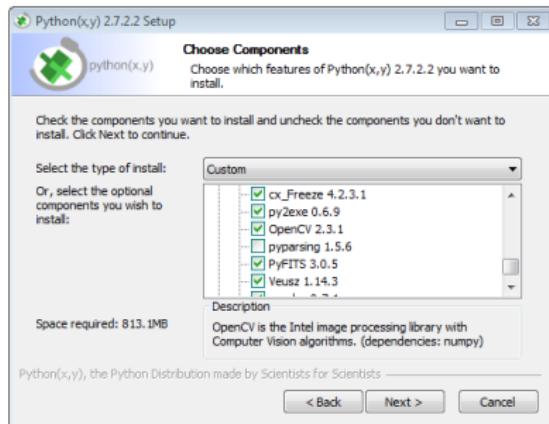


More info: <http://code.google.com/p/pythonxy/>



Installing on Windows – Python(x,y)

③ Select the components to install

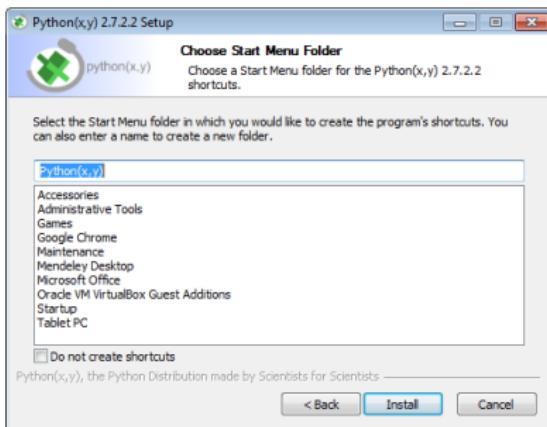
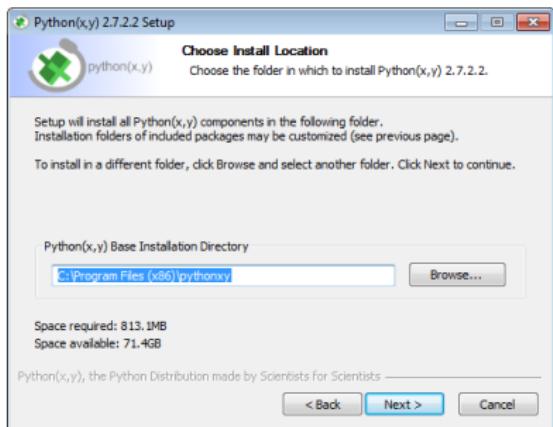


More info: <http://code.google.com/p/pythonxy/>

Installing on Windows – Python(x,y)



- ④ Select the installation location
- ⑤ Select the Start Menu folder
- ⑥ Press the **Install** button.

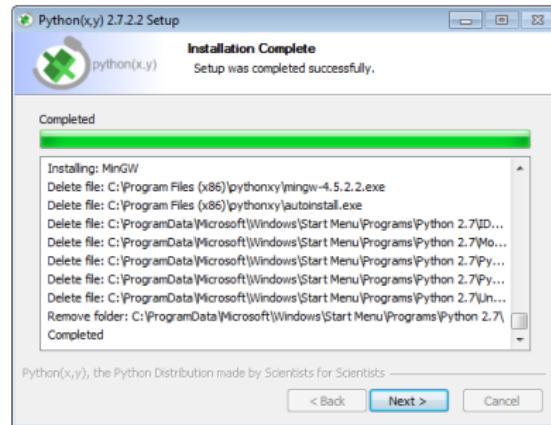
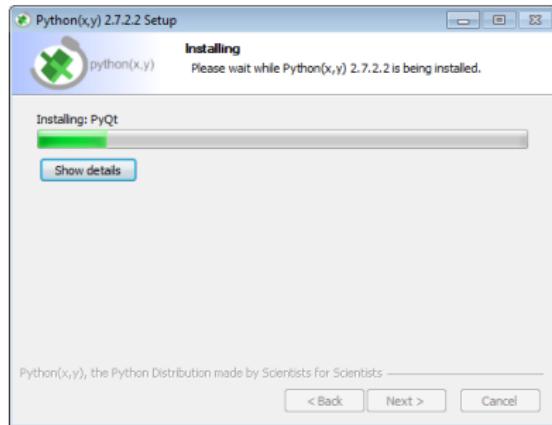


More info: <http://code.google.com/p/pythonxy/>

Installing on Windows – Python(x,y)



7 Wait for everything to install.



More info: <http://code.google.com/p/pythonxy/>

(Kitware / Google)

Python for MATLAB Users

June 2012

193 / 201

Installing on Windows – Python(x,y)



- ⑧ Press the **Finish** button and you are done!



More info: <http://code.google.com/p/pythonxy/>

Installing on OS X

This should get you setup with all the tools shown today in as “clean” of a way as possible.

Installing on OS X

- The good news:
 - OS X ships with Python installed
 - Homebrew is a great package manager for OS X (<http://mxcl.github.com/homebrew/>).
- The bad news:
 - We need to install xcode
 - Building everything can take a while

Installing on OS X

- First, we assume you have XCode installed.
- Next, install homebrew using the instructions here:
<https://github.com/mxcl/homebrew/wiki/installation>.
- Now that we have homebrew, we will install our binary dependencies

```
brew install gfortran # for scipy
brew install pyqt
brew install pyside
brew install opencv
```

- Now we're ready to setup the python stuff in a virtual environment.

Installing on OS X

- What is a virtual environment?
- Python virtual environments let you keep all of your dependencies in one place. (through environment variable magic)
- To install virtualenv, use the instruction here:
<http://pypi.python.org/pypi/virtualenv>
- Now create a virtual environment like so:

```
virtualenv down_with_matlab  
source down_with_matlab/bin/activate
```

- Now we are ready to install our python packages and applications

Installing on OS X

- The order and distinction between pip and easy_install is important.

```
easy_install readline
pip install sphinx
pip install spyder
pip install ipython
pip install tornado
easy_install zmq
pip install numpy
pip install scipy
pip install matplotlib
pip install sympy
easy_install SimpleITK
```

- Now you can use spyder, ipython, and all of these libraries as long as the virtual environment is activated.
- Bonus: None of this is cluttering up your systems site-packages.

Installing on Linux

- The same virtualenv, pip, and easy_install tools used for Mac OS X can be used for Linux, but there is an easier way.
- Most linux distributions have packages for scientific Python tools that make them easy to install.
 - Package are (usually) pre-built, so you don't have to wait for them to compile.
 - Dependencies are automatically installed. If you install Spyder, you automatically get Python, NumPy, SciPy, etc.
- A few toolkits may not be packaged and will still require pip or easy_install.

Installing on Linux - Ubuntu

- In Ubuntu, you can use the Ubuntu Software Center  to select and install packages.
- Equivalently, from the command line you can run

```
sudo apt-get install spyder
```

- Installing Spyder provides most of what is needed.
- Other useful packages to install are

```
sudo apt-get install ipython-notebook ipython-qtconsole python-opencv imagej  
sudo easy_install SimpleItk
```

- The apt-get commands will also work on other Debian-based OSes (like Debian and Mint).